

# Data Source Reliability & Trust Matrix

## HashInsight API Integration - Freshness, SLA, and Failover Strategy

**Version:** 1.0

**Date:** October 2025

**Document Type:** Technical Reference

## Executive Summary

HashInsight aggregates data from **7 independent sources** to ensure **99.95% data availability**. This document provides complete transparency on: - API update frequencies and latency - Caching strategies (TTL policies) - Failover and degradation mechanisms - Monitoring and alerting thresholds - Historical uptime and reliability metrics

**Key Metrics:** - **Data Freshness:** 1-60 seconds (depending on data type) - **Failover Time:** <2 seconds automatic fallback - **API Redundancy:** 2-3 sources per critical data point - **Cache Hit Rate:** 87% (production average)

## 1. Primary Data Sources

### 1.1 CoinGecko API

**Purpose:** Real-time cryptocurrency pricing

**Endpoint:** <https://api.coingecko.com/api/v3/simple/price>

Metric	Value	Details
<b>Update Frequency</b>	Every 30-60 seconds	CoinGecko updates BTC price every minute
<b>Our Polling Interval</b>	On-demand (cached)	Only fetch when cache expires
<b>Cache TTL</b>	60 seconds	Balance freshness vs API limits
<b>API Latency</b>	100-300ms	Typical response time (p95: 250ms)
<b>Rate Limit</b>	50 calls/minute (free tier)	Sufficient with caching
<b>Historical Uptime</b>	99.8% (2024)	Verified via <a href="https://status.coingecko.com">status.coingecko.com</a>
<b>Fallback</b>	Blockchain.info → Default \$80,000	2-tier fallback
<b>Monitoring Alert</b>	>500ms latency or >5 failures/min	PagerDuty alert

### Degradation Strategy:

```

# Step 1: Try CoinGecko (primary)
price = coingecko_api.get_btc_price()

# Step 2: Fallback to Blockchain.info
if price is None:
    price = blockchain_info_api.get_btc_price()

# Step 3: Use conservative default
if price is None:
    price = 80000.0 # Safe default for calculations
    logger.warning("All BTC price APIs failed, using default")

```

**Alert Conditions:** - API response time >500ms for 5 consecutive requests - Error rate >10% over 1-minute window - Price deviation >5% from previous value (data quality check)

## 1.2 Blockchain.info API

**Purpose:** Bitcoin network statistics (hashrate, difficulty) + Backup BTC price

**Endpoints:** - <https://blockchain.info/q/hashrate> - <https://blockchain.info/q/getdifficulty> - <https://blockchain.info/ticker>

Metric	Value	Details
<b>Update Frequency</b>	Network hashrate: ~10 minutes Difficulty: ~2 weeks Price: ~30 seconds	Based on Bitcoin network
<b>Our Polling Interval</b>	Hashrate: 5 min cache Difficulty: 10 min cache Price: 60 sec cache	Aligned with data change rate
<b>Cache TTL</b>	300s (hashrate), 600s (difficulty), 60s (price)	Optimized per data type
<b>API Latency</b>	80-200ms	Fast, simple API
<b>Rate Limit</b>	Unlimited (best effort)	No official limit documented
<b>Historical Uptime</b>	99.5% (2024)	Occasional brief outages
<b>Fallback</b>	Default values: 900 EH/s (hashrate), 119.12T (difficulty)	Industry average defaults
<b>Monitoring Alert</b>	>3 consecutive failures	Slack + PagerDuty

### Data Validation:

```
# Validate network hashrate (sanity check)
if hashrate_eh < 500 or hashrate_eh > 2000:
    logger.error(f"Invalid hashrate: {hashrate_eh} EH/s")
    return default_value # 900 EH/s

# Validate difficulty (must be positive)
if difficulty <= 0:
```

```
logger.error(f"Invalid difficulty: {difficulty}")
return 119.12e12 # Default
```

## 1.3 CoinWarz API

**Purpose:** Professional mining profitability data (multiple coins)

**Endpoint:** Configured via `COINWARZ_API_KEY` environment variable

Metric	Value	Details
<b>Update Frequency</b>	5-15 minutes	Mining pool data aggregation
<b>Our Polling Interval</b>	10 minutes (when used)	Matches their update cycle
<b>Cache TTL</b>	600 seconds (10 min)	Aligned with data freshness
<b>API Latency</b>	200-500ms	Multi-coin data retrieval
<b>Rate Limit</b>	1000 calls/day (API key)	Sufficient for our usage
<b>Historical Uptime</b>	98.5% (2024)	Occasional maintenance
<b>Fallback</b>	Calculate manually from BTC price + difficulty	Fallback calculation
<b>Monitoring Alert</b>	>1000ms latency or API key invalid	Email alert

**Usage Note:** CoinWarz is used for **multi-coin profitability comparisons** when users want to analyze non-BTC mining options.

## 1.4 Alternative.me API (Fear & Greed Index)

**Purpose:** Market sentiment indicator

**Endpoint:** <https://api.alternative.me/fng/>

Metric	Value	Details
<b>Update Frequency</b>	Daily (8am UTC)	Once per day update
<b>Our Polling Interval</b>	1 hour cache	Matches update frequency
<b>Cache TTL</b>	3600 seconds (1 hour)	No need for frequent polls
<b>API Latency</b>	100-250ms	Lightweight API
<b>Rate Limit</b>	No official limit	Public free API
<b>Historical Uptime</b>	99.2% (2024)	Generally reliable
<b>Fallback</b>	Omit from results (non-critical)	Optional indicator
<b>Monitoring Alert</b>	>5 consecutive failures	Low priority alert

**Degradation Strategy:** - If API fails: Simply exclude Fear & Greed Index from dashboard - Does NOT block mining calculations (non-critical data)

---

## 1.5 Ankr RPC (Bitcoin Blockchain)

**Purpose:** Real-time Bitcoin blockchain data

**Endpoint:** Free Bitcoin RPC service

Metric	Value	Details
<b>Update Frequency</b>	Real-time (per block ~10 min)	Live blockchain data
<b>Our Polling Interval</b>	As needed for Web3 features	Event-driven
<b>Cache TTL</b>	60 seconds (block height/hash)	Recent block data
<b>API Latency</b>	50-150ms	Fast RPC
<b>Rate Limit</b>	1500 req/sec (free tier)	More than sufficient
<b>Historical Uptime</b>	99.9% (2024)	Enterprise-grade infrastructure
<b>Fallback</b>	Direct node connection (if configured)	Backup RPC endpoint
<b>Monitoring Alert</b>	>500ms latency or block height stale >30min	Critical alert

**Primary Use Cases:** - SLA NFT smart contract interactions - Verifiable computing proof generation - Transparent hashrate validation

---

## 1.6 Derivatives Exchanges (Deribit, OKX, Binance)

**Purpose:** Funding rates, open interest, derivatives pressure indicators

**Endpoints:** Exchange-specific WebSocket + REST APIs

## Deribit API

Metric	Value	Details
<b>Update Frequency</b>	Real-time (WebSocket) REST: Every 1-5 seconds	Live derivatives data
<b>Our Polling Interval</b>	30 second cache (funding rate) Real-time (WebSocket for signals)	Mixed approach
<b>Cache TTL</b>	30 seconds (funding), 60s (open interest)	Balance freshness/ load
<b>API Latency</b>	30-100ms (WebSocket) 100-300ms (REST)	Low latency
<b>Rate Limit</b>	200 req/min (REST) Unlimited (WebSocket)	Generous limits
<b>Historical Uptime</b>	99.7% (2024)	Professional exchange
<b>Fallback</b>	OKX → Binance → Omit signal	3-tier fallback
<b>Monitoring Alert</b>	WebSocket disconnect >30s or >10 REST failures/min	High priority

## OKX API

Metric	Value	Details
<b>Update Frequency</b>	Real-time (WebSocket)	Live derivatives
<b>Cache TTL</b>	30 seconds	Short TTL for derivatives
<b>API Latency</b>	50-150ms	Fast exchange API
<b>Rate Limit</b>	100 req/sec	Sufficient
<b>Historical Uptime</b>	99.6% (2024)	Reliable
<b>Fallback</b>	Binance → Omit	Secondary fallback

## Binance API

Metric	Value	Details
<b>Update Frequency</b>	Real-time	Live data
<b>Cache TTL</b>	30 seconds	Derivatives freshness
<b>API Latency</b>	40-120ms	Very fast
<b>Rate Limit</b>	2400 req/min	High limit
<b>Historical Uptime</b>	99.8% (2024)	Top-tier reliability
<b>Fallback</b>	Last known value → Omit	Terminal fallback

### Signal Aggregation Strategy:

```
# Collect from all 3 exchanges
funding_rates = []
for exchange in [Deribit, OKX, Binance]:
    try:
        rate = exchange.get_funding_rate_cached()
        if rate is not None:
            funding_rates.append(rate)
```

```


except Exception as e:
    logger.warning(f"{exchange} failed: {e}")

# Use weighted average if ≥2 sources available
if len(funding_rates) >= 2:
    avg_funding = weighted_average(funding_rates)
else:
    # Insufficient data - omit from signals
    avg_funding = None

```

## 2. Caching Strategy Matrix

### 2.1 Multi-Tier Cache Architecture



**Cache Tiers:** 1. **L1 (In-Memory)**: Python dict, 1-60 second TTL, <1ms access 2. **L2 (Redis)**: Distributed cache, 60-3600 second TTL, 2-5ms access 3. **Request Coalescing**: Merge concurrent identical requests (9.8x improvement)

## 2.2 TTL Policy Table

Data Type	L1 TTL	L2 TTL	Reason
BTC Price	30s	60s	Balance freshness/API quota
Network Hashrate	2min	5min	Changes every ~10 minutes
Network Difficulty	5min	10min	Changes every 2 weeks
Fear & Greed Index	30min	1h	Updates once daily
Miner Specs	24h	7d	Static reference data
User Portfolio	1min	5min	User-specific, moderate freshness
Calculation Results	5min	15min	Reuse identical calculations
Funding Rates	15s	30s	Derivatives need freshness
API Keys (validation)	5min	15min	Security vs performance

## 2.3 Cache Invalidation Rules

**Automatic Invalidation:** - Time-based expiry (TTL)  
- Version-based (data schema changes)  
- Event-driven (e.g., new Bitcoin block → invalidate difficulty cache)


### Manual Invalidation (Admin):

```
# Flush specific cache key
curl -X POST https://api.hashinsight.io/admin/cache/flush \
-H "Authorization: Bearer <admin_token>" \
-d '{"key": "btc_price"}'

# Flush all caches (emergency)
curl -X POST https://api.hashinsight.io/admin/cache/flush_all
```

## 3. Failover & Degradation Strategies

### 3.1 Failover Decision Tree



### 3.2 Failover Tiers by Data Type

#### Critical Data (Must Have)

**BTC Price:** 1. CoinGecko (primary) 2. Blockchain.info (fallback 1) 3. Default \$80,000 (conservative fallback) 4. **Alert:** Engineering team immediately

**Network Hashrate:** 1. Blockchain.info (primary) 2. Default 900 EH/s (industry average) 3. **Alert:** Low priority (non-blocking)

**Network Difficulty:** 1. Blockchain.info (primary) 2. Default 119.12T (recent average) 3. **Alert:** Low priority

#### Optional Data (Nice to Have)

**Fear & Greed Index:** 1. Alternative.me (primary) 2. Omit from dashboard (graceful degradation) 3. **Alert:** No alert (optional feature)

**Derivatives Data:** 1. All 3 exchanges (aggregate) 2. Require  $\geq 2$  sources for signal validity 3. Omit signal if  $< 2$  sources available 4. **Alert:** Medium priority

### 3.3 Circuit Breaker Pattern

#### Implementation:

```
class CircuitBreaker:
    def __init__(self, failure_threshold=5, timeout=60):
        self.failure_count = 0
        self.failure_threshold = failure_threshold
        self.timeout = timeout
        self.last_failure_time = None
        self.state = 'CLOSED' # CLOSED, OPEN, HALF_OPEN

    def call(self, func):
        if self.state == 'OPEN':
            if time.time() - self.last_failure_time > self.timeout:
                self.state = 'HALF_OPEN'
            else:
                raise CircuitOpenError("Circuit breaker is OPEN")

        try:
            result = func()
            if self.state == 'HALF_OPEN':
                self.state = 'CLOSED'
                self.failure_count = 0
            return result
        except Exception as e:
            self.failure_count += 1
            self.last_failure_time = time.time()
            if self.failure_count >= self.failure_threshold:
                self.state = 'OPEN'
            raise
```

**Applied to APIs:** - CoinGecko: 5 failures → OPEN for 60 seconds → Auto-switch to Blockchain.info - Blockchain.info: 10 failures → OPEN for 120 seconds → Use defaults - Derivatives exchanges: 3 failures → OPEN for 30 seconds → Omit signal

## 4. Monitoring & Alerting

### 4.1 Real-Time Metrics Dashboard

#### Prometheus Metrics Exported:

```

# API call success rate (per source)
api_request_total{source="coingecko", status="success"} 1250
api_request_total{source="coingecko", status="failure"} 3

# API latency percentiles
api_latency_ms{source="coingecko", percentile="p95"} 218

# Cache hit rate
cache_hit_rate{tier="l1"} 0.65
cache_hit_rate{tier="l2"} 0.87

# Circuit breaker state
circuit_breaker_state{source="coingecko"} 0 # 0=CLOSED, 1=OPEN, 2=HALF_OPEN

```

**Grafana Dashboard Panels:** 1. **API Health Matrix** - Success rate per source (last 1h) 2. **Latency Heatmap** - p50/p95/p99 per source 3. **Cache Performance** - Hit rate, eviction rate 4. **Failover Events** - Timeline of fallback activations 5. **Data Freshness** - Age of cached data per key

## 4.2 Alert Thresholds

Alert	Condition	Severity	Channel	Action
<b>API Down</b>	>10 consecutive failures	Critical	PagerDuty + Slack	Immediate investigation
<b>High Latency</b>	p95 >500ms for 5 min	Warning	Slack	Monitor, may escalate
<b>Low Cache Hit Rate</b>	<60% for 10 min	Info	Slack	Optimize cache strategy
<b>Circuit Breaker Open</b>	State=OPEN for >5 min	Warning	Slack + Email	Check API status
<b>Stale Data</b>	Data age >2x TTL	Warning	Slack	API may be stuck
<b>Rate Limit Hit</b>	429 errors detected	Warning	Slack	Reduce poll frequency

## 4.3 Alert Runbook

### Runbook: CoinGecko API Failure

1. Check CoinGecko status page: <https://status.coingecko.com>
2. Verify failover to Blockchain.info is working (check logs)
3. If both fail, confirm default value (\$80,000) is being used
4. Estimate impact: Are calculations still running?
5. Communicate to users if prolonged (>30 min)
6. Escalate to CoinGecko support if SLA breach

### Runbook: All APIs Down (Nuclear Scenario)

1. Activate emergency mode: Use all default values
2. Display warning banner: "Data may be delayed"
3. Alert CTO immediately (critical incident)
4. Check network connectivity (DNS, firewall)
5. Consider CDN/proxy issue (Cloudflare, etc.)
6. Prepare incident report for post-mortem

## 5. Historical Reliability Data

---

### 5.1 Uptime Performance (2024)

API Source	Uptime (%)	MTBF	MTTR	Incidents
CoinGecko	99.8%	120h	8min	12
Blockchain.info	99.5%	96h	15min	18
CoinWarz	98.5%	48h	25min	32
Alternative.me	99.2%	80h	12min	15
Ankr RPC	99.9%	240h	3min	6
Deribit	99.7%	180h	5min	8
OKX	99.6%	150h	7min	10
Binance	99.8%	200h	4min	7

**Definitions:** - **MTBF:** Mean Time Between Failures - **MTTR:** Mean Time To Recovery -

**Incidents:** Documented outages >5 minutes

### 5.2 Data Quality Incidents

**2024 Q3 Analysis:** - **Price Anomalies:** 3 incidents (API returned stale data >30min old) -

**Network Hashrate Spikes:** 2 incidents (API bug, 2000+ EH/s reported) - **Missing Data:** 8 incidents (API returned null/error)

**Mitigations Implemented:** Data validation layer (sanity checks)

Anomaly detection (>5% deviation triggers review)

Dual-source validation for critical data

## 6. Data Freshness Guarantee

### 6.1 SLA Commitments

Data Type	Freshness SLA	Actual (p95)	Status
BTC Price	≤2 minutes	62 seconds	PASS
Network Hashrate	≤10 minutes	6.2 minutes	PASS
Network Difficulty	≤20 minutes	11 minutes	PASS
Funding Rates	≤5 minutes	1.8 minutes	PASS
Fear & Greed	≤2 hours	45 minutes	PASS

### 6.2 Freshness Monitoring

#### Real-Time Dashboard:

```
-- Query to check data age
SELECT
    data_type,
    MAX(updated_at) as last_update,
    NOW() - MAX(updated_at) as data_age
FROM market_data_cache
GROUP BY data_type
HAVING (NOW() - MAX(updated_at)) > interval '5 minutes'
```

**Automated Alerts:** - Data age >2x TTL → Slack warning - Data age >5x TTL → PagerDuty alert

## 7. Continuous Improvement

---

### 7.1 Weekly Review Metrics

**Every Monday @ 10am UTC:** - Review API uptime (target: 99.5% composite) - Analyze failover activations (minimize frequency) - Check cache hit rate (target: >80%) - Identify new data sources (reduce single-point failures)

### 7.2 Quarterly Initiatives

**Q4 2025 Roadmap:** 1. Add 3rd BTC price source (Coinbase Pro API) 2. Implement predictive cache warming (ML-based) 3. Upgrade to WebSocket for all exchange APIs 4. Deploy multi-region API proxies (reduce latency)

---

## 8. Appendix

---

### 8.1 Quick Reference: Default Values

```
DEFAULT_VALUES = {
    'btc_price': 80000.0,          # USD (conservative)
    'network_hashrate': 900.0,     # EH/s (2024 average)
    'network_difficulty': 119.12e12, # ~119T
    'fear_greed_index': None,      # Omit if unavailable
    'funding_rate': 0.0001,        # 0.01% (neutral)
}
```

### 8.2 Monitoring Commands

```
# Check API health
curl https://api.hashinsight.io/health

# View Prometheus metrics
curl https://api.hashinsight.io/metrics | grep api_request

# Force cache refresh (admin only)
```

```
curl -X POST https://api.hashinsight.io/admin/cache/refresh \  
-H "Authorization: Bearer $ADMIN_TOKEN"
```

## 8.3 Contact Information

**API Issues:** - **Email:** api-support@hashinsight.io - **Slack:** #api-monitoring - **PagerDuty:** On-call engineer (24/7)

**Data Quality Concerns:** - **Email:** data-team@hashinsight.io - **Jira:** Create ticket with tag  
`data-quality`

---

**Document Control:** - **Version:** 1.0 - **Last Updated:** October 3, 2025 - **Owner:** HashInsight Data Engineering Team - **Review Cycle:** Monthly - **Next Review:** November 1, 2025