

HashInsight

Operations Manual for HashInsight Platform

: v2.0
: 2025-10-03
: HashInsight Platform Operations Team
: INTERNAL - CONFIDENTIAL

- 1.
 - 2.
 - 3.
 - 4.
 - 5.
 - 6.
 - 7.
 - 8.
 - 9.
 - 10.
 - 11.
-

1

1.1

HashInsight

Web	Flask	3.0+	Python Web
WSGI	Gunicorn	21.0+	HTTP
	PostgreSQL	15+	(Neon)
	Redis /	7.0+	
	Web3.py + Base L2	-	
	Cryptography	41.0+	

1.2

```

HashInsight Platform
├── Core Application (main.py + app.py)
│   ├── Authentication & Authorization
│   ├── Session Management
│   └── Security Middleware
│
├── Mining Management Module
│   ├── Miner Dashboard
│   ├── Batch Calculator
│   └── Analytics Engine
│
├── CRM & Client Module
│   ├── Customer Management
│   ├── Billing System
│   └── Subscription Management
│
├── Blockchain Integration Module
│   ├── SLA NFT Management
│   ├── Verifiable Computing
│   └── Trust Reconciliation
│
└── Admin & Analytics Module
    ├── Market Data Analysis
    ├── Performance Monitoring
    └── Reporting System

```

1.3

- **KMS** : AWS KMS GCP KMS Azure Key Vault
- **mTLS** : CRL/OCSP
- **API** : `hs1_dev_key_*`
- **WireGuard VPN:**
- : SOC 2 / PCI DSS / GDPR

SLO

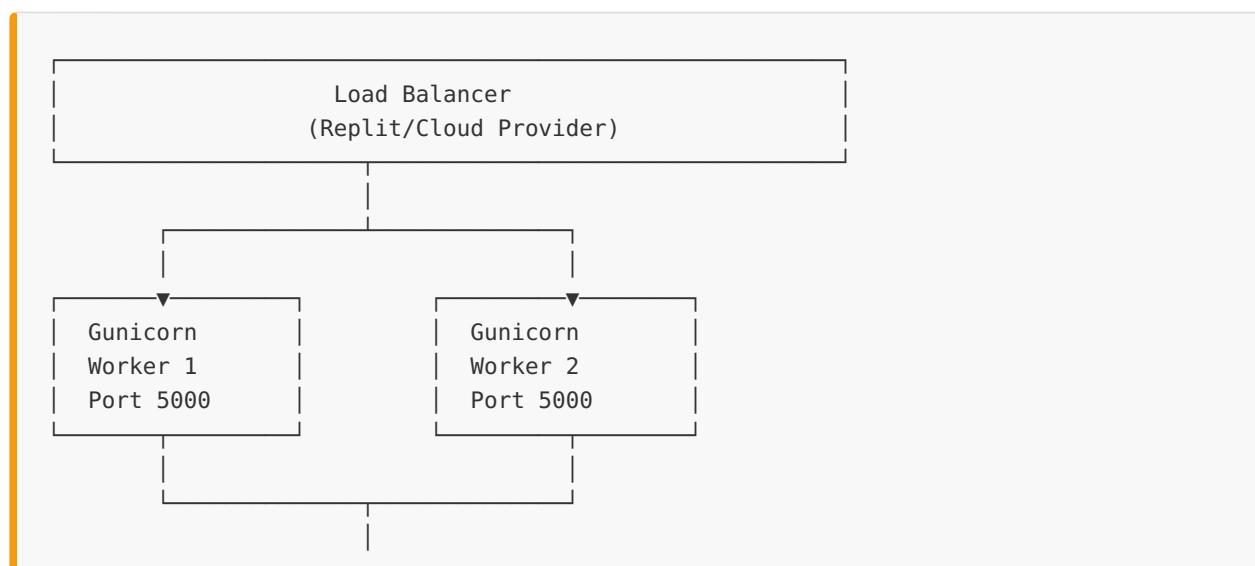
- : $\geq 99.95\%$ (≤ 21.6 /)
- : p95 ≤ 250 ms
- : $\leq 0.1\%$
- **Prometheus** :
- :

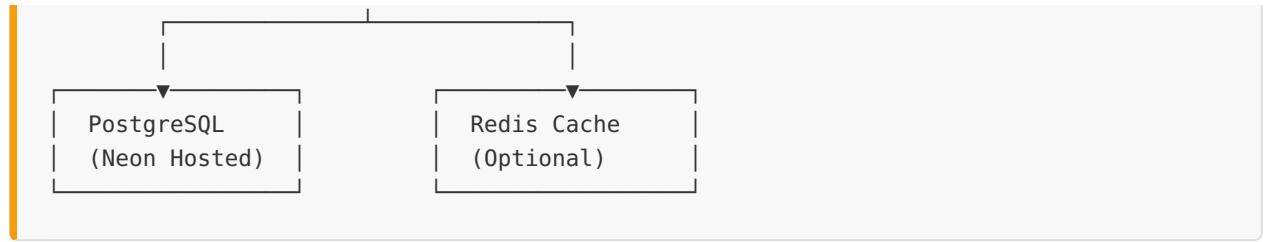
⚡

- **Request Coalescing:** 9.8

- : Redis +
- : PostgreSQL
- : 5000+

1.4





2

2.1

DATABASE_URL	String	PostgreSQL	postgresql://user:pass@host:5432/db
SESSION_SECRET	String	Flask (≥ 32)	your-secure-random-secret-key-here
ENCRYPTION_PASSWORD	String	(≥ 32)	encryption-master-key-32-chars-min

```
# .env
DATABASE_URL=postgresql://hashinsight_user:secure_password@neon-host.us-east-1.aws.neon.tech:5432/hash
SESSION_SECRET=generate_with_python_secrets_token_urlsafe_32
ENCRYPTION_PASSWORD=generate_with_python_secrets_token_urlsafe_32
```

```
# Python
import secrets
print(f"SESSION_SECRET={secrets.token_urlsafe(32)}")
print(f"ENCRYPTION_PASSWORD={secrets.token_urlsafe(32)}")
```

2.2

BLOCKCHAIN_ENABLED	Boolean	false	
BLOCKCHAIN_PRIVATE_KEY	String	-	(0x)
BLOCKCHAIN_NETWORK	String	base-sepolia	
BASE_RPC_URL	String	https://sepolia.base.org	Base L2 RPC

```
#  
BLOCKCHAIN_ENABLED=true  
BLOCKCHAIN_PRIVATE_KEY=0x1234567890abcdef...  
BLOCKCHAIN_NETWORK=base-sepolia
```

2.3

BACKUP_DIR	String	/tmp/backups	
BACKUP_ENCRYPTION_KEY	String	-	
BACKUP_RETENTION_DAYS	Integer	30	
BACKUP_STORAGE_TYPE	String	local	(s3/azure/gcs)
BACKUP_STORAGE_BUCKET	String	-	

```
# AWS S3  
BACKUP_DIR=/var/backups/hashinsight  
BACKUP_ENCRYPTION_KEY=backup-encryption-key-32-chars  
BACKUP_RETENTION_DAYS=30  
BACKUP_STORAGE_TYPE=s3  
BACKUP_STORAGE_BUCKET=hashinsight-backups  
BACKUP_STORAGE_REGION=us-east-1  
AWS_ACCESS_KEY_ID=AKIAXXXXXXXXXXXX  
AWS_SECRET_ACCESS_KEY=xxxxxxxxxxxx
```

2.4 KMS

AWS KMS

```
AWS_KMS_KEY_ID=arn:aws:kms:us-east-1:123456789:key/xxxxxx  
AWS_KMS_REGION=us-east-1  
AWS_ACCESS_KEY_ID=AKIAXXXXXXXX  
AWS_SECRET_ACCESS_KEY=xxxxxxxxxxxx
```

GCP KMS

```
GCP_KMS_PROJECT_ID=hashinsight-prod  
GCP_KMS_LOCATION=us-east1  
GCP_KMS_KEYRING=hashinsight-keyring  
GCP_KMS_KEY_ID=encryption-key  
GOOGLE_APPLICATION_CREDENTIALS=/path/to/service-account.json
```

Azure Key Vault

```
AZURE_KEY_VAULT_URL=https://hashinsight-vault.vault.azure.net/  
AZURE_TENANT_ID=xxxxxxxx-xxxx-xxxx-xxxx-xxxxxxxxxxxx  
AZURE_CLIENT_ID=xxxxxxxx-xxxx-xxxx-xxxx-xxxxxxxxxxxx  
AZURE_CLIENT_SECRET=xxxxxxxxxxxx
```

2.5

Configuration Settings			
ENABLE_BACKGROUND_SERVICES	Boolean	false	'1'
PROMETHEUS_PORT	Integer	9090	Prometheus
SLO_MEASUREMENT_WINDOW	Integer	30	SLO ()

: FAST_STARTUP

main.py

3.4

2.6 API

COINWARZ_API_KEY	CoinWarz API	https://www.coinwarz.com/api
COINGECKO_API_KEY	CoinGecko API	https://www.coingecko.com/api
SENDGRID_API_KEY	SendGrid	https://sendgrid.com

2.7

```
HashInsight/
└── config.py          #
└── .env               # (Git)
└── .env.example       #
└── replit.md          #
```

2.8

```
#  
python -c "  
import os  
required = ['DATABASE_URL', 'SESSION_SECRET', 'ENCRYPTION_PASSWORD']  
missing = [v for v in required if not os.getenv(v)]  
if missing:  
    print(f' Missing: {missing}')  
    exit(1)  
print(' All required variables set')  
"
```

3

3.1

()

```
gunicorn --bind 0.0.0.0:5000 --reuse-port --reload main:app
```

--bind 0.0.0.0:5000	5000 (Replit)
--reuse-port	worker
--reload	()
--workers 4	Worker (CPU ×2+1)
--timeout 120	Worker ()

```
gunicorn \
--bind 0.0.0.0:5000 \
--workers 4 \
--threads 2 \
--worker-class gthread \
--timeout 120 \
--max-requests 1000 \
--max-requests-jitter 50 \
--access-logfile - \
--error-logfile - \
--log-level info \
--preload \
main:app
```

3.2

△ : 5000

```
#  
lsof -i :5000  
  
#  
kill -9 $(lsof -t -i:5000)
```

3.3

```
#  
curl http://localhost:5000/health  
  
# ()  
curl http://localhost:5000/health/detailed  
  
#  
{  
    "status": "healthy",  
    "timestamp": "2025-10-03T12:00:00Z",  
    "database": "connected",  
    "cache": "available",  
    "version": "2.0.0"  
}
```

3.4 Fast Startup

HashInsight

main.py

config.py

```
# main.py  
fast_startup = os.environ.get("FAST_STARTUP", "1").lower() in ("1", "true", "yes") #  
skip_db_check = os.environ.get("SKIP_DATABASE_HEALTH_CHECK", "1").lower() in ("1", "true", "yes") #
```

config.py

```

# ()
export FAST_STARTUP=1
export SKIP_DATABASE_HEALTH_CHECK=1

# ()
export FAST_STARTUP=0
export SKIP_DATABASE_HEALTH_CHECK=0

```

⚠ - config.py - main.py - config.py

Fast Startup

1. (2-3) 2. 5 3. 4.

CI/CD

1. 2. 3. 10-15 4.

3.5

⚠ : ORM SQL

```

#
# app.py db.create_all()

#
python -c "from app import app, db; app.app_context().push(); db.create_all()"

```

3.6

```

# Step 1: Canary
# 1worker
gunicorn --bind 0.0.0.0:5001 --workers 1 main:app

# Step 2: Canary (5-10)
watch -n 5 'curl -s http://localhost:5001/health | jq'

# Step 3:
# : old(90%) -> new(10%)
#

# Step 4:
# old(0%) -> new(100%)

```

```
killall -9 gunicorn #  
gunicorn --bind 0.0.0.0:5000 --workers 4 main:app #
```

3.7

```
# SIGTERM ()  
kill -TERM $(cat /var/run/gunicorn.pid)  
  
# 30  
sleep 30  
  
# ()  
kill -KILL $(cat /var/run/gunicorn.pid)
```

3.8

```
#  
export LOG_LEVEL=INFO  
export LOG_FORMAT=json  
  
#  
tail -f /var/log/hashinsight/app.log  
  
#  
grep ERROR /var/log/hashinsight/app.log | tail -20  
  
# Replit  
# Replit Console
```

3.9

- [] (DATABASE_URL, SESSION_SECRET, ENCRYPTION_PASSWORD)
- []
- [] 5000
- [] 200
- [] (logs/audit.jsonl)
- []
- [] SSL (mTLS)
- [] KMS ()
- [] Prometheus (:9090/metrics)

4

4.1 SLO

HashInsight SLO

SLO

SLO			
	$\geq 99.95\%$	$\leq 21.6 /$	30
	$\geq 99.9\%$	$\leq 43.2 /$	30

SLO

SLO		
P50	$\leq 100\text{ms}$	5
P95	$\leq 250\text{ms}$	5
P99	$\leq 500\text{ms}$	5

SLO

SLO		
4xx	$\leq 1\%$	
5xx	$\leq 0.1\%$	

4.2 Prometheus

```
# monitoring/prometheus_exporter.py
from prometheus_client import Counter, Histogram, Gauge

#
request_count = Counter(
    'hashinsight_requests_total',
    'Total request count',
    ['method', 'endpoint', 'status']
)

#
request_latency = Histogram(
    'hashinsight_request_latency_seconds',
    'Request latency',
    ['method', 'endpoint'],
    buckets=[0.01, 0.05, 0.1, 0.25, 0.5, 1.0, 2.5, 5.0]
)

#
cache_hit_rate = Gauge(
    'hashinsight_cache_hit_rate',
    'Cache hit rate percentage'
)

#
db_query_duration = Histogram(
    'hashinsight_db_query_duration_seconds',
    'Database query duration',
    ['query_type'],
    buckets=[0.001, 0.01, 0.05, 0.1, 0.5, 1.0]
)

# SLO
slo_compliance = Gauge(
    'hashinsight_slo_compliance',
    'SLO compliance percentage',
    ['slo_type']
)
```

```
# Prometheus
curl http://localhost:9090/metrics

#
# HELP hashinsight_requests_total Total request count
```

```

# TYPE hashinsight_requests_total counter
hashinsight_requests_total{method="GET",endpoint="/dashboard",status="200"} 1234

# HELP hashinsight_request_latency_seconds Request latency
# TYPE hashinsight_request_latency_seconds histogram
hashinsight_request_latency_seconds_bucket{method="GET",endpoint="/api/miners",le="0.1"} 450
hashinsight_request_latency_seconds_bucket{method="GET",endpoint="/api/miners",le="0.25"} 480

```

4.3 Grafana

```
{
  "dashboard": {
    "title": "HashInsight Production Monitoring",
    "panels": [
      {
        "title": "Request Rate",
        "targets": [
          {
            "expr": "rate(hashinsight_requests_total[5m])"
          }
        ],
        "type": "timeseries"
      },
      {
        "title": "P95 Latency",
        "targets": [
          {
            "expr": "histogram_quantile(0.95, rate(hashinsight_request_latency_seconds_bucket[5m]))"
          }
        ],
        "type": "timeseries"
      },
      {
        "title": "Error Rate",
        "targets": [
          {
            "expr": "rate(hashinsight_requests_total{status=~\"5..\"}[5m]) / rate(hashinsight_requests_total)"
          }
        ],
        "type": "timeseries"
      },
      {
        "title": "SLO Compliance",
        "targets": [
          {
            "expr": "hashinsight_slo_compliance"
          }
        ],
        "type": "timeseries"
      }
    ]
  }
}
```

4.4

Prometheus

```
# prometheus/alerts.yml
groups:
  - name: hashinsight_alerts
    interval: 30s
    rules:
      #
      - alert: HighErrorRate
        expr: |
          rate(hashinsight_requests_total{status=~"5.."}[5m])
          / rate(hashinsight_requests_total[5m]) > 0.01
        for: 5m
        labels:
          severity: critical
        annotations:
          summary: "High 5xx error rate detected"
          description: "Error rate is {{ $value | humanizePercentage }}"

      #
      - alert: HighLatency
        expr: |
          histogram_quantile(0.95,
            rate(hashinsight_request_latency_seconds_bucket[5m])
            ) > 0.25
        for: 10m
        labels:
          severity: warning
        annotations:
          summary: "P95 latency exceeds SL0"
          description: "P95 latency is {{ $value }}s (SL0: 0.25s)"

      # SL0
      - alert: ErrorBudgetExhausted
        expr: hashinsight_slo_error_budget_remaining < 0.1
        for: 5m
        labels:
          severity: critical
        annotations:
          summary: "SL0 error budget nearly exhausted"
          description: "Only {{ $value | humanizePercentage }} budget remaining"

      #
      - alert: DatabaseConnectionFailure
        expr: hashinsight_db_connection_status == 0
        for: 2m
        labels:
          severity: critical
        annotations:
          summary: "Database connection failure"
```

```

        description: "Unable to connect to PostgreSQL"

#
- alert: LowCacheHitRate
  expr: hashinsight_cache_hit_rate < 50
  for: 15m
  labels:
    severity: warning
  annotations:
    summary: "Low cache hit rate"
    description: "Cache hit rate is {{ $value }}%"

```

4.5

HashInsight

```

# monitoring/circuit_breaker.py
from monitoring.circuit_breaker import CircuitBreaker, circuit_breaker

#
db_breaker = CircuitBreaker(
    failure_threshold=5,      # 5
    recovery_timeout=60,      # 60
    name="database_queries"
)

# API
@circuit_breaker(
    failure_threshold=3,
    recovery_timeout=30,
    name="external_api"
)
def call_external_api():
    response = requests.get("https://api.coinwarz.com/")
    return response.json()

```

```

#
curl http://localhost:5000/api/circuit-breakers

#
{
  "database_queries": {
    "state": "closed",
    "failure_count": 0,
    "total_calls": 1234,
    "success_rate": "99.8%"
  },

```

```
"external_api": {  
    "state": "half_open",  
    "failure_count": 3,  
    "total_calls": 456,  
    "success_rate": "95.2%"  
}  
}
```

4.6

Slack

```
#  
SLACK_WEBHOOK_URL=https://hooks.slack.com/services/T00/B00/yyyyy  
ALERT_SLACK_CHANNEL=#hashinsight-alerts
```

PagerDuty

```
PAGERDUTY_API_KEY=xxxxx  
PAGERDUTY_SERVICE_KEY=xxxxx
```

4.7

- [] Prometheus ([:9090/targets](#))
- [] Grafana
- []
- [] Slack/PagerDuty
- [] SLO
- []

5

5.1

HashInsight [backup/backup_manager.py](#)

- **PostgreSQL** (pg_dump)
- **AES-256** ()
- **gzip** ()
- (S3/Azure/GCS)
- (SHA256)

```
# cron
# /etc/cron.d/hashinsight-backup

# 2
0 2 * * * /usr/bin/python3 /app/backup/backup_manager.py --type full

# 4
0 */4 * * * /usr/bin/python3 /app/backup/backup_manager.py --type incremental

# 3
0 3 * * 0 /usr/bin/python3 /app/backup/backup_manager.py --upload
```

5.2

```
#
python backup/backup_manager.py

#
: hashinsight_backup_20251003_140000.sql.gz.enc
: 245.3 MB
: AES-256
: e3b0c44298fc1c149afbf4c8996fb92427ae41e4649b934ca495991b7852b855
: 45.2s
```

5.3

	30	
	7	4
	12	
	12	1

```
# 30
python backup/backup_manager.py --cleanup --days 30

#
: hashinsight_backup_20250903_*.sql.gz.enc
: 2.1 GB
```

5.4 RTO/RPO

RTO ()	≤ 4	~ 2
RPO ()	≤ 15	~ 4

5.5

Step 1:

```
#
ls -lh /tmp/backups/
#
# (S3)
aws s3 ls s3://hashinsight-backups/
```

Step 2: ()

```
# S3
aws s3 cp s3://hashinsight-backups/hashinsight_backup_20251003_020000.sql.gz.enc /tmp/restore/

# Azure
az storage blob download \
--account-name hashinsight \
--container-name backups \
--name hashinsight_backup_20251003_020000.sql.gz.enc \
--file /tmp/restore/backup.sql.gz.enc
```

Step 3:

```
# backup_manager
python backup/backup_manager.py \
--decrypt /tmp/restore/hashinsight_backup_20251003_020000.sql.gz.enc \
--output /tmp/restore/backup.sql.gz

# ()
openssl enc -d -aes-256-cbc \
-in hashinsight_backup_20251003_020000.sql.gz.enc \
-out backup.sql.gz \
-pass env:BACKUP_ENCRYPTION_KEY
```

Step 4:

```
gunzip /tmp/restore/backup.sql.gz
# : backup.sql
```

Step 5:

```
# △ : !
#
# PostgreSQL
psql $DATABASE_URL < /tmp/restore/backup.sql

#
psql $DATABASE_URL -c "SELECT COUNT(*) FROM users;"
psql $DATABASE_URL -c "SELECT COUNT(*) FROM miners;"
```

Step 6:

```
#  
systemctl restart hashinsight  
  
#  
curl http://localhost:5000/health  
  
#  
curl http://localhost:5000/api/miners | jq '.count'
```

5.6

:

1.

```
#  
createdb hashinsight_dr_test  
export DATABASE_URL=postgresql://localhost/hashinsight_dr_test
```

1.

```
#  
psql $DATABASE_URL -c "DROP SCHEMA public CASCADE; CREATE SCHEMA public;"
```

1. (5.5)

2.

```
#  
python -c "  
from app import app, db  
from models import User, Miner  
with app.app_context():  
    assert db.session.query(User).count() > 0  
    assert db.session.query(Miner).count() > 0  
    print(' ')"
```

1. - - - - -

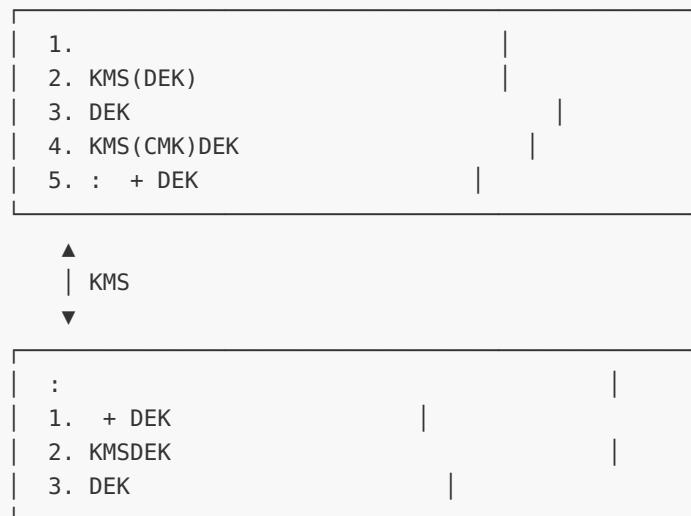
5.7

```
#  
stat -c '%y' /tmp/backups/hashinsight_backup_*.sql.gz.enc | tail -1  
  
#  
du -h /tmp/backups/hashinsight_backup_*.sql.gz.enc  
  
#  
python backup/backup_manager.py --verify /tmp/backups/hashinsight_backup_20251003_020000.sql.gz.enc
```

6

6.1 KMS

HashInsight KMS ([common/crypto/envelope.py](#))



KMS

AWS KMS

```
# AWS KMS  
from common.crypto.envelope import KMSClient, KMSProvider, EncryptionContext
```

```

kms_config = {
    'key_id': 'arn:aws:kms:us-east-1:123456789:key/xxxxx',
    'region': 'us-east-1'
}

client = KMSClient(KMSProvider.AWS_KMS, kms_config)

#
context = EncryptionContext(
    purpose="user_data_encryption",
    tenant_id="tenant_123"
)

ciphertext = client.encrypt_secret(
    plaintext="sensitive data",
    key_id=kms_config['key_id'],
    context=context
)

```

GCP KMS

```

# GCP KMS
kms_config = {
    'project_id': 'hashinsight-prod',
    'location': 'us-east1',
    'keyring': 'hashinsight-keyring',
    'key_id': 'encryption-key'
}

client = KMSClient(KMSProvider.GCP_KMS, kms_config)

```

Azure Key Vault

```

# Azure Key Vault
kms_config = {
    'vault_url': 'https://hashinsight-vault.vault.azure.net/',
    'key_name': 'encryption-key'
}

client = KMSClient(KMSProvider.AZURE_KEY_VAULT, kms_config)

```

```

# 1. KMS
aws kms create-key --description "HashInsight Master Key v2"

# 2.

```

```

export AWS_KMS_KEY_ID=arn:aws:kms:us-east-1:123456789:key/new-key-id

# 3. ()
python scripts/rotate_encryption_keys.py --old-key OLD_KEY_ID --new-key NEW_KEY_ID

# 4.
python scripts/verify_encryption.py --key-id NEW_KEY_ID

# 5. (90)
aws kms disable-key --key-id OLD_KEY_ID

```

6.2 mTLS

```

# 1. CA
openssl genrsa -out ca.key 4096
openssl req -new -x509 -days 3650 -key ca.key -out ca.crt \
-subj "/C=US/O=HashInsight/CN=HashInsight Root CA"

# 2.
openssl genrsa -out server.key 4096
openssl req -new -key server.key -out server.csr \
-subj "/C=US/O=HashInsight/CN=*.hashinsight.net"
openssl x509 -req -days 365 -in server.csr -CA ca.crt -CAkey ca.key \
-CAcreateserial -out server.crt

# 3.
openssl genrsa -out client.key 4096
openssl req -new -key client.key -out client.csr \
-subj "/C=US/O=HashInsight/CN=client.hashinsight.net"
openssl x509 -req -days 365 -in client.csr -CA ca.crt -CAkey ca.key \
-CAcreateserial -out client.crt

```

mTLS

```

#
export MTLS_ENABLED=true
export MTLS_CA_CERT_PATH=/app/certs/ca.crt
export MTLS_SERVER_CERT_PATH=/app/certs/server.crt
export MTLS_SERVER_KEY_PATH=/app/certs/server.key
export MTLS_VERIFY_CLIENT=true
export MTLS_ALLOWED_DN_PATTERNS="CN=*.hashinsight.net,O=HashInsight,C=US"

```

mTLS

```
from common.mtls_auth import require_mtls

@app.route('/api/admin/sensitive')
@require_mtls()
def sensitive_endpoint():
    #
    client_dn = g.client_cert_subject
    return jsonify({"message": f"Authenticated as {client_dn}"})
```

6.3 API

HashInsight API : `hsı_{env}_key_{random}`

: - : `hsı_prod_key_a1b2c3d4e5f6g7h8` - : `hsı_dev_key_x9y8z7w6v5u4t3s2`

API

```
#  
python scripts/create_api_key.py \  
--user-id 123 \  
--permissions "miners:read,miners:write" \  
--expires-in 90  
  
#  
API  
: hsi_prod_key_a1b2c3d4e5f6g7h8  
ID: 123  
: miners:read,miners:write  
: 2025-12-31T23:59:59Z  
▲
```

API

```
# 1.  
new_key=$(python scripts/create_api_key.py --user-id 123 --copy-from old_key_id)  
  
# 2. (7)  
#  
  
# 3.  
curl -H "Authorization: Bearer $new_key" http://localhost:5000/api/miners
```

```
# 4.  
python scripts/revoke_api_key.py --key-id old_key_id
```

6.4 WireGuard

Hub

```
# 1. Hub  
sudo bash wireguard/hub_setup.sh  
  
# 2.  
sudo ufw allow 51820/udp  
sudo ufw enable  
  
# 3. WireGuard  
sudo systemctl enable wg-quick@wg0  
sudo systemctl start wg-quick@wg0  
  
# 4.  
sudo wg show
```

```
# 1.  
cd wireguard/site-gateway  
python key_manager.py --generate-keys --site beijing-dc1  
  
# 2. Hub  
sudo nano /etc/wireguard/wg0.conf  
  
# Peer  
[Peer]  
PublicKey = site_public_key  
AllowedIPs = 10.8.1.0/24  
Endpoint = beijing-gateway.hashinsight.net:51820  
  
# 3.  
sudo wg-quick down wg0  
sudo wg-quick up wg0
```

6.5

HashInsight ([audit/audit_logger.py](#))

- (/ /)
- (CRUD)
-
-
-
- API
-

```
{
  "timestamp": "2025-10-03T12:00:00.000Z",
  "event_id": "a1b2c3d4e5f6g7h8",
  "level": "INFO",
  "category": "authentication",
  "action": "login",
  "user_id": "123",
  "user_email": "user@example.com",
  "user_role": "admin",
  "ip_address": "192.168.1.100",
  "status": "success",
  "details": {
    "login_method": "password",
    "two_factor": true
  }
}
```

```
# 100
tail -100 logs/audit.jsonl

#
jq 'select(.user_email == "admin@hashinsight.net")' logs/audit.jsonl

#
jq 'select(.action == "login_failed")' logs/audit.jsonl

# 24
jq --arg date "$(date -d '24 hours ago' -Iseconds)" \
'select(.timestamp > $date and .level == "SECURITY")' \
logs/audit.jsonl
```

6.6

HashInsight

SOC 2 Type II

-
- (TLS 1.3)
- (AES-256)
- ()
-
-

PCI DSS ()

-
- (KMS)
- (WireGuard)
-

GDPR

-
-
- ()
-

6.7

- [] KMS (/)
- [] mTLS ()
- [] API (90)
- [] SSL/TLS
- []
- []
- [] (SSL)
- [] ≥ 256

- [] ()
 - [] ()
-

7

7.1

```
$ gunicorn main:app  
[ERROR] Application failed to start
```

1.

```
# 5000  
lsof -i :5000  
  
#  
kill -9 $(lsof -t -i:5000)
```

2.

```
#  
python3 << EOF  
import os  
required = ['DATABASE_URL', 'SESSION_SECRET', 'ENCRYPTION_PASSWORD']  
for var in required:  
    value = os.getenv(var)  
    if not value:  
        print(f" Missing: {var}")  
    else:  
        print(f" {var}: {'*' * 8} (set)")  
EOF
```

3.

```
# PostgreSQL  
psql $DATABASE_URL -c "SELECT version();"
```

```
# Neon
# https://console.neon.tech
```

4. Python

```
#
python3 -c "
import flask
import gunicorn
import psycopg2
import sqlalchemy
print(' All dependencies OK')
"
```

5.

```
#
export FLASK_DEBUG=1
python3 main.py

#
```

ModuleNotFoundError: No module named 'flask'		pip install -r requirements.txt
OperationalError: could not connect to server		DATABASE_URL Neon
ValueError: SECRET_KEY must be set	SESSION_SECRET	export SESSION_SECRET=\$(python -c 'import secrets; print(secrets.token_urlsafe(32))')
Address already in use		kill -9 \$(lsof -t -i:5000)

7.2

```
sqlalchemy.exc.OperationalError: (psycopg2.OperationalError)
could not connect to server: Connection timed out
```

1.

```
# DATABASE_URL
python3 << 'EOF'
import os
from urllib.parse import urlparse
url = os.getenv('DATABASE_URL')
parsed = urlparse(url)
print(f"Host: {parsed.hostname}")
print(f"Port: {parsed.port}")
print(f"Database: {parsed.path[1:]}")
print(f"User: {parsed.username}")
EOF
```

2.

```
# hostport
export DB_HOST=$(python3 -c "from urllib.parse import urlparse; import os; print(urlparse(os.getenv('DATABASE_URL')).hostname)")
export DB_PORT=$(python3 -c "from urllib.parse import urlparse; import os; print(urlparse(os.getenv('DATABASE_URL')).port)")

# TCP
nc -zv $DB_HOST $DB_PORT
```

3.

```
# (config.py)
SQLALCHEMY_ENGINE_OPTIONS = {
    'pool_size': 10,          # 10
    'pool_recycle': 300,      # 5
    'pool_pre_ping': True,   #
    'pool_timeout': 30,       # 30
    'max_overflow': 20,       # 20
    'connect_args': {
        'connect_timeout': 15 # 15
    }
}
```

4. Neon

```
# Neon
# https://console.neon.tech/app/projects

# :
# - Active () -
# - Idle () -
# - Suspended () -
```

5.

```
# app.py
from sqlalchemy import event, exc
from sqlalchemy.pool import Pool

@event.listens_for(Pool, "connect")
def receive_connect(dbapi_conn, connection_record):
    connection_record.info['pid'] = os.getpid()

@event.listens_for(Pool, "checkout")
def receive_checkout(dbapi_conn, connection_record, connection_proxy):
    pid = os.getpid()
    if connection_record.info['pid'] != pid:
        connection_record.connection = connection_proxy.connection = None
        raise exc.DisconnectionError(
            "Connection record belongs to pid %s, "
            "attempting to check out in pid %s" %
            (connection_record.info['pid'], pid)
        )
```

7.3

-
- 50%
- Redis

1. Redis

```
# Redis
redis-cli ping
# : PONG
```

```
# Redis
redis-cli info memory

#
redis-cli dbsize
```

2.

```
# API
curl http://localhost:5000/api/cache/stats

#
{
    "cache_type": "redis",
    "hit_rate": 75.5,
    "total_requests": 10000,
    "hits": 7550,
    "misses": 2450
}
```

3.

```
# cache_manager.py
if redis_available:
    cache_backend = RedisCache()
else:
    logger.warning("Redis unavailable, falling back to memory cache")
    cache_backend = MemoryCache()
```

4.

```
#
redis-cli FLUSHDB

#
redis-cli --scan --pattern 'hashinsight:' | xargs redis-cli DEL
```

7.4

- API p95 > 250ms
-

- CPU/

1. Request Coalescing

```
# curl http://localhost:5000/api/performance/coalescing-stats

#
{
  "enabled": true,
  "performance_improvement": "9.8x",
  "deduplicated_requests": 5432,
  "total_requests": 53210
}
```

2.

```
-- (PostgreSQL)
ALTER DATABASE hashinsight_db SET log_min_duration_statement = 1000;

-- 
SELECT query, calls, total_time, mean_time
FROM pg_stat_statements
WHERE mean_time > 1000
ORDER BY mean_time DESC
LIMIT 10;
```

3.

```
# python database/optimize_indexes.py --analyze

#
python database/optimize_indexes.py --suggest
```

4.

```
# CPU
top -b -n 1 | grep gunicorn

#
ps aux | grep gunicorn | awk '{sum+=$6} END {print sum/1024 " MB"}'
```

```
#  
psql $DATABASE_URL -c "SELECT count(*) FROM pg_stat_activity WHERE datname = 'hashinsight_db';"
```

5.

```
#  
from werkzeug.middleware.profiler import ProfilerMiddleware  
  
app.wsgi_app = ProfilerMiddleware(  
    app.wsgi_app,  
    restrictions=[10],  
    profile_dir='./profiles'  
)
```

7.5

```
ERROR:blockchain_integration: : ENCRYPTION_PASSWORD  
ERROR:sla_nft_routes: SLA
```

1. ENCRYPTION_PASSWORD

```
#  
echo $ENCRYPTION_PASSWORD  
  
#  
export ENCRYPTION_PASSWORD=$(python3 -c 'import secrets; print(secrets.token_urlsafe(32))')
```

2.

```
#  
python3 << 'EOF'  
import os  
config = {  
    'BLOCKCHAIN_ENABLED': os.getenv('BLOCKCHAIN_ENABLED'),  
    'BLOCKCHAIN_PRIVATE_KEY': '****' if os.getenv('BLOCKCHAIN_PRIVATE_KEY') else None,  
    'BLOCKCHAIN_NETWORK': os.getenv('BLOCKCHAIN_NETWORK', 'base-sepolia'),  
    'BASE_RPC_URL': os.getenv('BASE_RPC_URL', 'https://sepolia.base.org')  
}  
for k, v in config.items():  
    status = '' if v else ''
```

```
    print(f"{status} {k}: {v}")
EOF
```

3. Web3

```
python3 << 'EOF'
from web3 import Web3
import os

rpc_url = os.getenv('BASE_RPC_URL', 'https://sepolia.base.org')
w3 = Web3(Web3.HTTPProvider(rpc_url))

if w3.is_connected():
    print(f" Web3 connected to {rpc_url}")
    print(f"Block number: {w3.eth.block_number}")
else:
    print(f" Web3 connection failed")
EOF
```

4.

```
# 0x64
python3 << 'EOF'
import os
import re

private_key = os.getenv('BLOCKCHAIN_PRIVATE_KEY', '')
if re.match(r'^0x[0-9a-fA-F]{64}$', private_key):
    print(" Private key format valid")
else:
    print(" Invalid private key format")
    print("Expected: 0x + 64 hex characters")
EOF
```

7.6

	(stdout)	
	logs/audit.jsonl	JSON Lines
	(stderr)	
Gunicorn	/var/log/gunicorn/	
PostgreSQL	Neon Console	
Workflow	/tmp/logs/	

```
# tail -f /var/log/hashinsight/app.log  
  
# grep ERROR /var/log/hashinsight/app.log | tail -50  
  
# tail -f logs/audit.jsonl | jq '.'  
  
# Gunicorn  
tail -f /var/log/gunicorn/access.log
```

8

8.1

8.1.1

```
# python database/optimize_indexes.py --auto
```

```

#
python database/optimize_indexes.py --analyze --suggest

# :
# : 45
# ...
# :
#   - CREATE INDEX idx_miners_user_id ON miners(user_id)
#   - CREATE INDEX idx_market_data_timestamp ON market_analytics(created_at DESC)

```

8.1.2

```

-- 
SELECT
    pid,
    username,
    application_name,
    client_addr,
    state,
    query_start
FROM pg_stat_activity
WHERE datname = 'hashinsight_db'
ORDER BY query_start;

-- 
SELECT pg_terminate_backend(pid)
FROM pg_stat_activity
WHERE datname = 'hashinsight_db'
    AND state = 'idle'
    AND state_change < NOW() - INTERVAL '30 minutes';

```

8.1.3

```

-- pg_stat_statements
CREATE EXTENSION IF NOT EXISTS pg_stat_statements;

-- Top 10
SELECT
    substring(query, 1, 100) AS short_query,
    calls,
    total_time,
    mean_time,
    max_time
FROM pg_stat_statements
WHERE query NOT LIKE '%pg_stat_statements%'
ORDER BY mean_time DESC
LIMIT 10;

```

```
--  
SELECT pg_stat_statements_reset();
```

8.1.4

```
# (90)  
python scripts/cleanup_old_data.py --days 90  
  
# :  
# 90market_analytics: 12,543  
# 90audit: 45,123  
# : 2.3 GB
```

8.2

8.2.1 Redis

```
# HashInsight  
redis-cli --scan --pattern 'hashinsight:' | xargs redis-cli DEL  
  
#  
redis-cli --scan --pattern 'hashinsight:miners:' | xargs redis-cli DEL  
  
# (Redis)  
redis-cli --scan --pattern 'hashinsight:' | while read key; do  
    redis-cli TTL "$key"  
done
```

8.2.2

```
# API  
curl http://localhost:5000/api/cache/memory-stats  
  
#  
{  
    "cache_type": "memory",  
    "size_mb": 124.5,  
    "entries": 5432,  
    "hit_rate": 82.3,  
    "evictions": 234  
}
```

8.2.3 Request Coalescer

```
#  
curl http://localhost:5000/api/performance/coalescing-stats | jq  
  
#  
curl -X POST http://localhost:5000/api/performance/clear-coalescing-cache
```

8.3

8.3.1

```
# CSV (5000)  
# : name,model,hashrate_th,power_w,efficiency  
  
#  
python batch/batch_import_manager.py \  
--file miners_upload_5000.csv \  
--user-id 123 \  
--validate  
  
# :  
#  
# : 5000  
# * ...  
# [██████████] 100%  
# : 5000 (: 45.2s)
```

8.3.2

HashInsight 15

```
#  
python modules/analytics/engines/analytics_engine.py --collect-now  
  
#  
curl http://localhost:5000/api/analytics/collection-status  
  
#  
{  
    "last_collection": "2025-10-03T12:15:00Z",  
    "next_collection": "2025-10-03T12:30:00Z",  
    "status": "healthy",  
    "data_points_today": 8  
}
```

8.3.3

```
# (config.py)
ANALYTICS_COLLECTION_INTERVAL = 15 #
ANALYTICS_MAX_DATA_POINTS_PER_DAY = 10 #

# / 0
export ENABLE_BACKGROUND_SERVICES=1 #
export ENABLE_BACKGROUND_SERVICES=0 #
```

8.4

8.4.1

```
#
# : http://localhost:5000/admin/users/create

#
python scripts/create_user.py \
--email admin@hashinsight.net \
--username admin \
--role owner \
--password-prompt

#
# :
# : *****
#
# ID: 123
# Email: admin@hashinsight.net
# Role: owner
```

8.4.2

owner		
admin		
broker		
client		

```
#  
python scripts/change_user_role.py \  
--user-id 123 \  
--new-role admin  
  
#  
python scripts/bulk_import_users.py --file users.csv
```

8.4.3

```
#  
from decorators import has_permission  
  
@app.route('/api/miners/delete/<int:miner_id>', methods=['DELETE'])  
@login_required  
@requires_permission('miners:delete')  
def delete_miner(miner_id):  
    # miners:delete  
    pass
```

8.5

8.5.1

HashInsight

```
# config.py  
ANALYTICS_MAX_DATA_POINTS_PER_DAY = 10 # 10
```

8.5.2

```
# 90  
python scripts/archive_historical_data.py --days 90 --storage s3  
  
# :  
# ...  
# market_analytics: 123,456  
# calculation_history: 45,678  
# ↑ s3://hashinsight-archive/...  
#  
# : 5.2 GB
```

8.5.3

```
# 6
python scripts/archive_audit_logs.py --months 6

# ()
# audit/audit_logger.py 100MB
```

9

9.1 Request Coalescing

HashInsight Request Coalescing **9.8**

```
:
Request 1 → API Call → Response 1
Request 2 → API Call → Response 2
Request 3 → API Call → Response 3
(3API)

Request Coalescing:
Request 1 ↗
Request 2 ↗ Single API Call → Response → Shared Result
Request 3 ↗
(1API67%)
```

```
# cache_manager.py
REQUEST_COALESCING_ENABLED = True
COALESCING_TIMEOUT = 100 # 100ms
COALESCING_MAX_WAIT = 500 # 500ms
```

```
# Request Coalescing
curl http://localhost:5000/api/performance/coalescing-stats | jq
```

```
#  
{  
    "enabled": true,  
    "total_requests": 98234,  
    "deduplicated_requests": 85432,  
    "api_calls_saved": 85432,  
    "performance_improvement": "9.8x",  
    "average_wait_time_ms": 45  
}
```

9.2

9.2.1

```
--  
CREATE INDEX idx_miners_user_id ON miners(user_id);  
CREATE INDEX idx_market_data_timestamp ON market_analytics(created_at DESC);  
CREATE INDEX idx_calculations_user_created ON calculations(user_id, created_at);  
  
-- ()  
CREATE INDEX idx_miners_user_model ON miners(user_id, model);  
  
-- ()  
CREATE INDEX idx_active_miners ON miners(user_id) WHERE status = 'active';
```

9.2.2

EXPLAIN

```
EXPLAIN ANALYZE  
SELECT m.*, u.username  
FROM miners m  
JOIN users u ON m.user_id = u.id  
WHERE m.status = 'active'  
ORDER BY m.created_at DESC  
LIMIT 100;
```

```
--
```

N+1

```
# - N+1  
miners = Miner.query.filter_by(user_id=user_id).all()  
for miner in miners:  
    print(miner.user.username) #
```

```

#     - JOIN
miners = Miner.query.join(User).filter(Miner.user_id == user_id).all()
for miner in miners:
    print(miner.user.username)  #

```

9.2.3

```

# config.py
SQLALCHEMY_ENGINE_OPTIONS = {
    'pool_size': 10,           # (worker × 2)
    'pool_recycle': 300,       # 5
    'pool_pre_ping': True,    # (Neon)
    'pool_timeout': 30,        # 30
    'max_overflow': 20,        # 20
    'connect_args': {
        'connect_timeout': 15,      #
        'application_name': 'hashinsight', #
        'options': '-c statement_timeout=30000' # 30
    }
}

```

9.3

9.3.1

```

#
L1:  ()
    ↓ miss
L2: Redis ()
    ↓ miss
L3:  ()

#
@cache_manager.cached(ttl=300, level='L1') # 5
def get_user_profile(user_id):
    return db.session.query(User).get(user_id)

@cache_manager.cached(ttl=3600, level='L2') # 1
def get_market_data():
    return fetch_from_api()

```

9.3.2 TTL

TTL
5
5
1
10
24

9.3.3

```
#  
python scripts/warmup_cache.py  
  
#  
# 1.  
# 2.  
# 3.  
# 4.  
  
# :  
#   : 1,234  
#   : +  
#   : 10  
#   : 12.3s
```

9.4

9.4.1

```
# -  
for miner in miners:  
    daily_revenue = calculate_revenue(  
        miner.hashrate,  
        btc_price,  
        network_difficulty  
    )  
    # : 5000 × 10ms = 50
```

```

# -
import numpy as np

hashrates = np.array([m.hashrate for m in miners])
revenues = calculate_revenue_vectorized(
    hashrates,
    btc_price,
    network_difficulty
)
# : ~500ms (100)

```

9.4.2

```

from concurrent.futures import ThreadPoolExecutor

#
def process_miner_batch(miners, batch_size=100):
    with ThreadPoolExecutor(max_workers=8) as executor:
        futures = []
        for i in range(0, len(miners), batch_size):
            batch = miners[i:i+batch_size]
            future = executor.submit(process_batch, batch)
            futures.append(future)

    results = [f.result() for f in futures]
    return results

```

9.4.3

```

# (ORM)
from sqlalchemy import insert

#
for data in large_dataset:
    db.session.add(Model(**data))
    db.session.commit() #

#
db.session.bulk_insert_mappings(Model, large_dataset)
db.session.commit() #

# - bulk insert
stmt = insert(Model).values(large_dataset)
db.session.execute(stmt)
db.session.commit()

```

9.5

```
#  
python scripts/performance_benchmark.py  
  
# :  
# API:  
#   - /api/miners: p50=45ms, p95=120ms, p99=250ms  
#   - /api/dashboard: p50=80ms, p95=200ms, p99=400ms  
# :  
#   - : 15ms  
#   - (>1s): 0  
# :  
#   - : 78.5%  
#   - Request Coalescing: 9.8x
```

10

10.1 On-Call

09:00-18:00	DevOps		
18:00-09:00			On-Call Manager
/			On-Call Manager

- **PagerDuty:**
- **Slack:** #incident-response
- **Grafana:**
- **Incident.io:**

10.2

P0	15	30		
P1	30	1	API >5s >5%	
P2	2	4		
P3	4	8		
P4	1	2	UI bug	

10.3

P0/P1

1. [0-5]
 - |→ PagerDuty
 - |→ Slack
 - |→
2. [5-15]
 - |→ (ACK)
 - |→
 - |→
 - |→
3. [15-30]
 - |→
 - |→
 - |→
 - |→
4. [30+]
 - |→
 - |→
 - |→
 - |→

10.4

1:

```
sqlalchemy.exc.TimeoutError: QueuePool limit of size 10 overflow 20 reached
```

```
# 1.  
psql $DATABASE_URL -c "  
SELECT pg_terminate_backend(pid)  
FROM pg_stat_activity  
WHERE state = 'idle'  
    AND state_change < NOW() - INTERVAL '10 minutes';  
"  
  
# 2. ()  
export SQLALCHEMY_POOL_SIZE=20  
export SQLALCHEMY_MAX_OVERFLOW=40  
  
# 3.  
systemctl restart hashinsight  
  
# 4.  
watch -n 5 "psql $DATABASE_URL -c 'SELECT count(*) FROM pg_stat_activity;'"
```

2: OOM

```
MemoryError: Unable to allocate array  
Killed (OOM)
```

```
# 1. worker  
kill -HUP $(cat /var/run/gunicorn.pid)  
  
# 2.  
redis-cli FLUSHDB  
  
# 3. worker ()  
gunicorn --workers 2 --max-requests 500 main:app
```

```
# 4.  
watch -n 5 'free -h'
```

3:

```
- - 401/403 - IP
```

```
# 1.  
export RATE_LIMIT_ENABLED=true  
export RATE_LIMIT_REQUESTS_PER_MINUTE=10  
  
# 2. IP ()  
ufw deny from 192.168.1.0/24  
  
# 3. Cloudflare ()  
# Cloudflare Dashboard -> Security -> DDoS  
  
# 4.  
tail -1000 /var/log/gunicorn/access.log | \  
awk '{print $1}' | sort | uniq -c | sort -rn | head -20
```

10.5

```
# 1.  
git log --oneline -10  
  
# 2.  
git revert HEAD --no-edit  
#  
git checkout v1.9.5  
  
# 3.  
git push origin main  
  
# 4.  
curl http://localhost:5000/health | jq '.version'  
  
# 5.  
# Slack: "v1.9.5"
```

```
# △ : !  
  
# 1.  
python backup/backup_manager.py --type emergency  
  
# 2.  
psql $DATABASE_URL < /tmp/backups/hashinsight_backup_20251003_020000.sql  
  
# 3.  
python scripts/verify_database_integrity.py  
  
# 4.  
systemctl restart hashinsight
```

10.6

```
[P0 ] HashInsight  
:  
: 2025-10-03 14:23 UTC  
:  
:  
30  
:  
: #incident-2025-10-03-db-outage  
: @john.doe
```

```
[] HashInsight  
:  
: 2025-10-03 14:45 UTC  
:  
:  
:  
: 15:00 UTC
```

```
[] HashInsight  
:  
: 2025-10-03 15:12 UTC  
: 49  
: PostgreSQL  
:
```

24

10.7 (Postmortem)

```
# HashInsight - 2025-10-03  
  
##  
- ****: INC-2025-1003-001  
- ****: P0  

```

```
- [ ]  
##
```

A:

```
#  
gunicorn --bind 0.0.0.0:5000 --workers 4 main:app  
  
#  
kill -HUP $(cat /var/run/gunicorn.pid)  
  
#  
kill -TERM $(cat /var/run/gunicorn.pid)  
  
#  
ps aux | grep gunicorn  
  
#  
tail -f /var/log/hashinsight/app.log
```

```
#  
psql $DATABASE_URL  
  
#  
\dt  
  
#  
SELECT count(*) FROM pg_stat_activity;  
  
#  
pg_dump $DATABASE_URL > backup.sql
```

```
#  
psql $DATABASE_URL < backup.sql
```

```
# Redis  
redis-cli  
  
#  
redis-cli KEYS 'hashinsight:'*'  
  
#  
redis-cli FLUSHDB  
  
#  
redis-cli INFO memory
```

```
#  
curl http://localhost:5000/health | jq  
  
# Prometheus  
curl http://localhost:9090/metrics  
  
# SLO  
curl http://localhost:5000/api/slo/status | jq
```

B:

.env

```
#  
DATABASE_URL=postgresql://user:pass@host:5432/hashinsight_db  
SESSION_SECRET=your-secret-key-min-32-chars  
ENCRYPTION_PASSWORD=encryption-key-min-32-chars  
  
#  
BLOCKCHAIN_ENABLED=true  
BLOCKCHAIN_PRIVATE_KEY=0x1234567890abcdef...  
BLOCKCHAIN_NETWORK=base-sepolia  
BASE_RPC_URL=https://sepolia.base.org  
  
#  
BACKUP_DIR=/var/backups/hashinsight  
BACKUP_ENCRYPTION_KEY=backup-encryption-key
```

```

BACKUP_RETENTION_DAYS=30
BACKUP_STORAGE_TYPE=s3
BACKUP_STORAGE_BUCKET=hashinsight-backups

# KMS (AWS)
AWS_KMS_KEY_ID=arn:aws:kms:us-east-1:123456789:key/xxxxxx
AWS_KMS_REGION=us-east-1
AWS_ACCESS_KEY_ID=AKIAXXXXXXXXX
AWS_SECRET_ACCESS_KEY=xxxxxxxxxxxx

#
ENABLE_BACKGROUND_SERVICES=0 # 1
PROMETHEUS_PORT=9090
SLO_MEASUREMENT_WINDOW=30

# API
COINWARZ_API_KEY=your-coinwarz-api-key
COINGECKO_API_KEY=your-coingecko-api-key

```

systemd

```

# /etc/systemd/system/hashinsight.service

[Unit]
Description=HashInsight Platform
After=network.target postgresql.service

[Service]
Type=notify
User=hashinsight
Group=hashinsight
WorkingDirectory=/opt/hashinsight
EnvironmentFile=/opt/hashinsight/.env

ExecStart=/opt/hashinsight/venv/bin/gunicorn \
--bind 0.0.0.0:5000 \
--workers 4 \
--threads 2 \
--worker-class gthread \
--timeout 120 \
--max-requests 1000 \
--access-logfile /var/log/hashinsight/access.log \
--error-logfile /var/log/hashinsight/error.log \
--pid /var/run/gunicorn.pid \
main:app

ExecReload=/bin/kill -s HUP $MAINPID
KillMode=mixed
KillSignal=SIGTERM
TimeoutStopSec=30

Restart=on-failure

```

```
RestartSec=10  
[Install]  
WantedBy=multi-user.target
```

C:

Prometheus

hashinsight_requests_total	Counter	
hashinsight_request_latency_seconds	Histogram	
hashinsight_error_rate	Gauge	
hashinsight_cache_hit_rate	Gauge	
hashinsight_db_query_duration_seconds	Histogram	
hashinsight_db_connection_pool_size	Gauge	
hashinsight_db_connection_pool_active	Gauge	
hashinsight_slo_compliance	Gauge	SLO
hashinsight_slo_error_budget_remaining	Gauge	
hashinsight_circuit_breaker_state	Gauge	

SLO

SLO			
	99.95%	<99.9%	<99.5%
P95	≤250ms	>200ms	>300ms
	≤0.1%	>0.5%	>1%
	21.6min/	<20%	<10%

D:

	HTTP		
AUTH_001	401		
AUTH_002	403		
AUTH_003	401	Session	
DB_001	500		DATABASE_URL
DB_002	500		
CACHE_001	503	Redis	Redis
API_001	429		
API_002	500	API	
BLOCKCHAIN_001	500		BLOCKCHAIN_PRIVATE_KEY
ENCRYPTION_001	500		ENCRYPTION_PASSWORD

E:

System Configuration Log			
v2.0	2025-10-03	HashInsight Ops Team	
v1.9	2025-09-15	DevOps Team	Request Coalescing
v1.8	2025-08-20	Security Team	KMS mTLS
v1.7	2025-07-10	Platform Team	SLO

: HashInsight Platform Operations Team
:
: ops@hashinsight.net
: <https://github.com/hashinsight/operations-manual>
: 2025-10-03
: 2026-01-03
