

# HashInsight Platform Operations Manual

**Document Version:** v2.0

**Last Updated:** October 3, 2025

**Maintained By:** HashInsight Platform Operations Team

**Classification:** INTERNAL - CONFIDENTIAL

## Table of Contents

- 
- [1. System Architecture Overview](#)
  - [2. Environment Configuration](#)
  - [3. Deployment Operations Guide](#)
  - [4. Monitoring & Alerting](#)
  - [5. Backup & Recovery](#)
  - [6. Security Operations](#)
  - [7. Troubleshooting Guide](#)
  - [8. Daily Operations](#)
  - [9. Performance Tuning Guide](#)
  - [10. Emergency Response](#)
  - [11. Appendices](#)
- 

## Chapter 1: System Architecture Overview

### 1.1 Technology Stack

HashInsight is an enterprise-grade Bitcoin mining management platform built with modern technology:

Component	Technology	Version	Purpose
<b>Web Framework</b>	Flask	3.0+	Python web application framework
<b>WSGI Server</b>	Gunicorn	21.0+	Production-grade HTTP server
<b>Database</b>	PostgreSQL	15+	Primary data storage (Neon hosted)
<b>Cache</b>	Redis / Memory Cache	7.0+	Performance optimization layer
<b>Blockchain</b>	Web3.py + Base L2	-	Blockchain integration
<b>Encryption</b>	Cryptography	41.0+	Enterprise-grade encryption

## 1.2 Module Architecture

The system uses a **fully page-isolated architecture** with independently deployed modules:

```

HashInsight Platform
├── Core Application (main.py + app.py)
│   ├── Authentication & Authorization
│   ├── Session Management
│   └── Security Middleware
|
├── Mining Management Module
│   ├── Miner Dashboard
│   ├── Batch Calculator
│   └── Analytics Engine
|
├── CRM & Client Module
│   ├── Customer Management
│   ├── Billing System
│   └── Subscription Management
|
├── Blockchain Integration Module
│   ├── SLA NFT Management
│   ├── Verifiable Computing
│   └── Trust Reconciliation
|
└── Admin & Analytics Module
    └── Market Data Analysis

```

```
  ├── Performance Monitoring  
  └── Reporting System
```

## 1.3 Enterprise Upgrade Achievements

### Security Enhancements

- **KMS Key Management:** Support for AWS KMS, GCP KMS, Azure Key Vault
- **mTLS Mutual Authentication:** Client certificate validation, CRL/OCSP checking
- **API Key System:** Secure keys based on `hs1_dev_key_*` format
- **WireGuard VPN:** Enterprise-grade network isolation
- **Audit Logging:** SOC 2 / PCI DSS / GDPR compliant

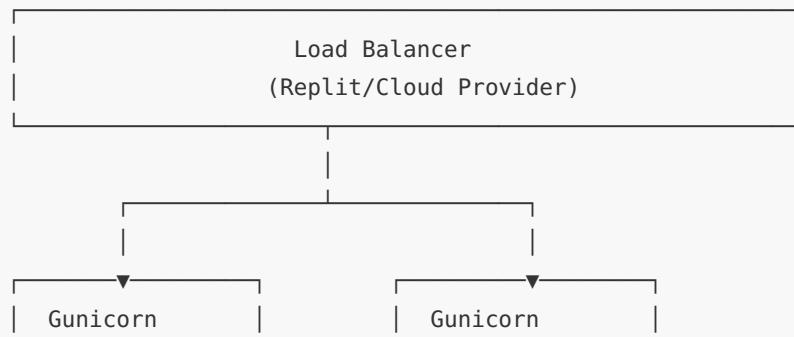
### SLO Monitoring

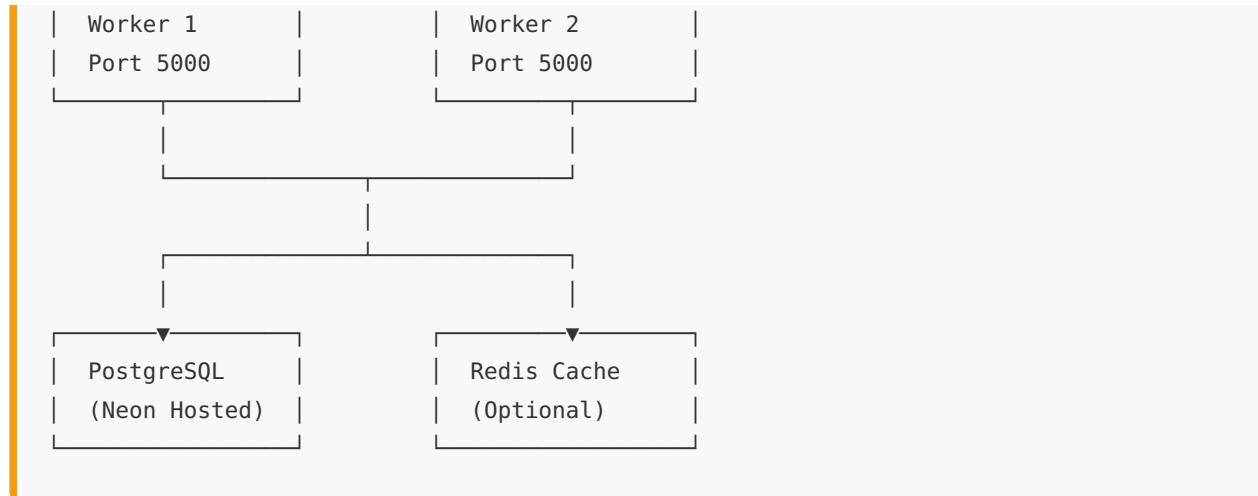
- **Availability:**  $\geq 99.95\%$  (error budget  $\leq 21.6$  minutes/month)
- **Latency:** p95  $\leq 250\text{ms}$
- **Error Rate:**  $\leq 0.1\%$
- **Prometheus Metrics:** Comprehensive performance monitoring
- **Circuit Breakers:** Prevent cascading failures

### ⚡ Performance Optimization

- **Request Coalescing:** 9.8x performance improvement
- **Multi-tier Caching:** Redis + memory cache
- **Connection Pooling:** PostgreSQL connection optimization
- **Batch Processing:** Support for 5000+ concurrent miner imports

## 1.4 Deployment Topology





## Chapter 2: Environment Configuration

### 2.1 Required Environment Variables

These environment variables **must** be set, or the system will fail to start:

Variable Name	Type	Description	Example
DATABASE_URL	String	PostgreSQL connection string	postgresql://user:pass@host:5432/db
SESSION_SECRET	String	Flask session secret ( $\geq 32$ chars)	your-secure-random-secret-key-here
ENCRYPTION_PASSWORD	String	Data encryption master key ( $\geq 32$ chars)	encryption-master-key-32-chars-min

#### Configuration Example

```

# .env file example
DATABASE_URL=postgresql://hashinsight_user:secure_password@neon-host.us-east-1.aws.neon.tech:5432/hash
SESSION_SECRET=generate_with_python_secrets_token_urlsafe_32
ENCRYPTION_PASSWORD=generate_with_python_secrets_token_urlsafe_32

```

## Generate Secure Keys

```
# Generate secure random keys using Python
import secrets
print(f"SESSION_SECRET={secrets.token_urlsafe(32)}")
print(f"ENCRYPTION_PASSWORD={secrets.token_urlsafe(32)}")
```

## 2.2 Blockchain Integration Configuration

Variable Name	Type	Default	Description
BLOCKCHAIN_ENABLED	Boolean	false	Enable blockchain features
BLOCKCHAIN_PRIVATE_KEY	String	-	Ethereum private key (0x prefix)
BLOCKCHAIN_NETWORK	String	base-sepolia	Blockchain network
BASE_RPC_URL	String	<a href="https://sepolia.base.org">https://sepolia.base.org</a>	Base L2 RPC endpoint

```
# Blockchain configuration example
BLOCKCHAIN_ENABLED=true
BLOCKCHAIN_PRIVATE_KEY=0x1234567890abcdef...
BLOCKCHAIN_NETWORK=base-sepolia
```

## 2.3 Backup System Configuration

Variable Name	Type	Default	Description
BACKUP_DIR	String	/tmp/backups	Backup storage directory
BACKUP_ENCRYPTION_KEY	String	-	Backup encryption key
BACKUP_RETENTION_DAYS	Integer	30	Backup retention period (days)
BACKUP_STORAGE_TYPE	String	local	Remote storage type (s3/azure/gcs)
BACKUP_STORAGE_BUCKET	String	-	Remote storage bucket name

```
# AWS S3 backup configuration
BACKUP_DIR=/var/backups/hashinsight
BACKUP_ENCRYPTION_KEY=backup-encryption-key-32-chars
BACKUP_RETENTION_DAYS=30
BACKUP_STORAGE_TYPE=s3
BACKUP_STORAGE_BUCKET=hashinsight-backups
BACKUP_STORAGE_REGION=us-east-1
AWS_ACCESS_KEY_ID=AKIAXXXXXXXX
AWS_SECRET_ACCESS_KEY=xxxxxxxxxxxx
```

## 2.4 KMS Key Management Configuration

### AWS KMS

```
AWS_KMS_KEY_ID=arn:aws:kms:us-east-1:123456789:key/xxxxxx
AWS_KMS_REGION=us-east-1
AWS_ACCESS_KEY_ID=AKIAXXXXXXXX
AWS_SECRET_ACCESS_KEY=xxxxxxxxxxxx
```

### GCP KMS

```
GCP_KMS_PROJECT_ID=hashinsight-prod
GCP_KMS_LOCATION=us-east1
GCP_KMS_KEYRING=hashinsight-keyring
```

```
GCP_KMS_KEY_ID=encryption-key  
GOOGLE_APPLICATION_CREDENTIALS=/path/to/service-account.json
```

## Azure Key Vault

```
AZURE_KEY_VAULT_URL=https://hashinsight-vault.vault.azure.net/  
AZURE_TENANT_ID=xxxxxxxx-xxxx-xxxx-xxxx-xxxxxxxxxxxx  
AZURE_CLIENT_ID=xxxxxxxx-xxxx-xxxx-xxxx-xxxxxxxxxxxx  
AZURE_CLIENT_SECRET=xxxxxxxxxxxx
```

## 2.5 Monitoring and Performance Configuration

Variable Name	Type	Default	Description
ENABLE_BACKGROUND_SERVICES	Boolean	false	Enable background data collection (must explicitly set to '1' to enable)
PROMETHEUS_PORT	Integer	9090	Prometheus export port
SLO_MEASUREMENT_WINDOW	Integer	30	SLO measurement window (minutes)

**Note:** FAST\_STARTUP mode is hardcoded in `main.py`, not configured through environment variables. See Section 3.4.

## 2.6 External API Configuration

Variable Name	Description	How to Obtain
COINWARZ_API_KEY	CoinWarz mining data API	<a href="https://www.coinwarz.com/api">https://www.coinwarz.com/api</a>
COINGECKO_API_KEY	CoinGecko price API	<a href="https://www.coingecko.com/api">https://www.coingecko.com/api</a>
SENDGRID_API_KEY	SendGrid email service	<a href="https://sendgrid.com">https://sendgrid.com</a>

## 2.7 Configuration File Locations

```
HashInsight/  
└── config.py # Main configuration file (single source of truth)
```

```
|── .env                      # Environment variables (not committed to Git)
|── .env.example               # Environment variable template
└── replit.md                 # System architecture documentation
```

## 2.8 Configuration Validation

```
# Check required environment variables
python -c "
import os
required = ['DATABASE_URL', 'SESSION_SECRET', 'ENCRYPTION_PASSWORD']
missing = [v for v in required if not os.getenv(v)]
if missing:
    print(f'  Missing: {missing}')
    exit(1)
print('  All required variables set')
"
```

# Chapter 3: Deployment Operations Guide

---

## 3.1 Startup Commands

### Standard Startup (Production Environment)

```
gunicorn --bind 0.0.0.0:5000 --reuse-port --reload main:app
```

## Parameter Explanation

Parameter	Description
--bind 0.0.0.0:5000	Bind to all network interfaces, port 5000 (Replit required)
--reuse-port	Allow multiple workers to bind to same port
--reload	Auto-reload on code changes (development)
--workers 4	Number of worker processes (CPU cores × 2 + 1)
--timeout 120	Worker timeout (seconds)

## Complete Production Startup

```
gunicorn \
--bind 0.0.0.0:5000 \
--workers 4 \
--threads 2 \
--worker-class gthread \
--timeout 120 \
--max-requests 1000 \
--max-requests-jitter 50 \
--access-logfile - \
--error-logfile - \
--log-level info \
--preload \
main:app
```

## 3.2 Port Configuration

⚠ **Important:** Frontend application **must** bind to port 5000

```
# Check port usage
lsof -i :5000

# Force kill occupying process
kill -9 $(lsof -t -i:5000)
```

### 3.3 Health Checks

The system provides health check endpoints for monitoring:

```
# Basic health check
curl http://localhost:5000/health

# Detailed health check (includes database status)
curl http://localhost:5000/health/detailed

# Expected response
{
    "status": "healthy",
    "timestamp": "2025-10-03T12:00:00Z",
    "database": "connected",
    "cache": "available",
    "version": "2.0.0"
}
```

### 3.4 Fast Startup Mode

HashInsight's fast startup mode is **hardcoded** in `main.py` as a built-in behavior, not managed through `config.py`.

#### Implementation Location

```
# Hardcoded implementation in main.py
fast_startup = os.environ.get("FAST_STARTUP", "1").lower() in ("1", "true", "yes") # Default: enabled
skip_db_check = os.environ.get("SKIP_DATABASE_HEALTH_CHECK", "1").lower() in ("1", "true", "yes") # Disabled by default
```

#### Control Method (Optional Environment Variables)

Although these variables are not defined in `config.py`, they can be temporarily adjusted via environment variables:

```
# Enable fast startup (default behavior, no need to set)
export FAST_STARTUP=1
export SKIP_DATABASE_HEALTH_CHECK=1

# Disable fast startup (full initialization)
```

```
export FAST_STARTUP=0  
export SKIP_DATABASE_HEALTH_CHECK=0
```

⚠ **Important Note:** - These variables are **not in the `config.py`** **single source of truth** -  
They are temporary control switches for the `main.py` startup script - Production  
configuration should be managed through `config.py`

## Fast Startup Behavior

When enabled (default): 1. Main application starts immediately (2-3 seconds) 2. Background services delayed 5 seconds (if `ENABLE_BACKGROUND_SERVICES=1`) 3. Database health checks skipped 4. Suitable for CI/CD rapid deployment

When disabled: 1. Full database health check 2. Synchronous startup of all services 3. Longer startup time (10-15 seconds) 4. Suitable for initial production deployment

## 3.5 Database Migration

⚠ **Important:** System uses ORM auto-migration, **avoid manual SQL operations**

```
# Database tables created automatically on application startup  
# See db.create_all() in app.py  
  
# Manually trigger migration if needed  
python -c "from app import app, db; app.app_context().push(); db.create_all()"
```

## 3.6 Rolling Deployment Strategy

### Canary Deployment Process

```
# Step 1: Deploy new version to Canary instance  
# Only 1 worker runs new version  
gunicorn --bind 0.0.0.0:5001 --workers 1 main:app  
  
# Step 2: Monitor Canary instance (5-10 minutes)  
watch -n 5 'curl -s http://localhost:5001/health | jq'  
  
# Step 3: Gradually shift traffic  
# Use load balancer to adjust weights: old(90%) -> new(10%)  
# Monitor error rate and latency  
  
# Step 4: Full cutover
```

```
# old(0%) -> new(100%)
killall -9 gunicorn # Stop old version
gunicorn --bind 0.0.0.0:5000 --workers 4 main:app # Start new version
```

## 3.7 Graceful Shutdown

```
# Send SIGTERM signal (graceful shutdown)
kill -TERM $(cat /var/run/gunicorn.pid)

# Wait 30 seconds to process existing requests
sleep 30

# Force stop (if still running)
kill -KILL $(cat /var/run/gunicorn.pid)
```

## 3.8 Log Management

```
# Enable structured logging at startup
export LOG_LEVEL=INFO
export LOG_FORMAT=json

# View real-time logs
tail -f /var/log/hashinsight/app.log

# View error logs
grep ERROR /var/log/hashinsight/app.log | tail -20

# Replit environment logs
# Logs output to stdout, view through Replit Console
```

## 3.9 Deployment Checklist

- [ ] Environment variables configured (DATABASE\_URL, SESSION\_SECRET, ENCRYPTION\_PASSWORD)
- [ ] Database connection working
- [ ] Port 5000 available
- [ ] Health check endpoint returns 200
- [ ] Audit log directory writable ( `logs/audit.jsonl` )
- [ ] Backup directory created

- [ ] SSL certificates valid (if mTLS enabled)
  - [ ] KMS keys accessible (if enabled)
  - [ ] Prometheus metrics accessible ( :9090/metrics )
- 

## Chapter 4: Monitoring & Alerting

### 4.1 SLO Definitions

HashInsight follows strict SLO standards:

#### Availability SLO

Metric	Target	Error Budget	Measurement Period
Availability	≥99.95%	≤21.6 minutes/month	30-day rolling
Success Rate	≥99.9%	≤43.2 minutes/month	30-day rolling

#### Latency SLO

Percentile	Target	Measurement Window
P50	≤100ms	5 minutes
P95	≤250ms	5 minutes
P99	≤500ms	5 minutes

#### Error Rate SLO

Type	Target	Threshold
4xx errors	≤1%	Warning
5xx errors	≤0.1%	Critical

## 4.2 Prometheus Metrics

### System Metrics Export

```
# monitoring/prometheus_exporter.py
from prometheus_client import Counter, Histogram, Gauge

# Request count
request_count = Counter(
    'hashinsight_requests_total',
    'Total request count',
    ['method', 'endpoint', 'status']
)

# Request latency
request_latency = Histogram(
    'hashinsight_request_latency_seconds',
    'Request latency',
    ['method', 'endpoint'],
    buckets=[0.01, 0.05, 0.1, 0.25, 0.5, 1.0, 2.5, 5.0]
)

# Cache hit rate
cache_hit_rate = Gauge(
    'hashinsight_cache_hit_rate',
    'Cache hit rate percentage'
)

# Database query duration
db_query_duration = Histogram(
    'hashinsight_db_query_duration_seconds',
    'Database query duration',
    ['query_type'],
    buckets=[0.001, 0.01, 0.05, 0.1, 0.5, 1.0]
)

# SLO compliance
slo_compliance = Gauge(
    'hashinsight_slo_compliance',
    'SLO compliance percentage',
    ['slo_type']
)
```

## Metrics Access

```
# View Prometheus metrics
curl http://localhost:9090/metrics

# Example output
# HELP hashinsight_requests_total Total request count
# TYPE hashinsight_requests_total counter
hashinsight_requests_total{method="GET",endpoint="/dashboard",status="200"} 1234

# HELP hashinsight_request_latency_seconds Request latency
# TYPE hashinsight_request_latency_seconds histogram
hashinsight_request_latency_seconds_bucket{method="GET",endpoint="/api/miners",le="0.1"} 450
hashinsight_request_latency_seconds_bucket{method="GET",endpoint="/api/miners",le="0.25"} 480
```

## 4.3 Grafana Dashboard

### Core Monitoring Panels

```
{
  "dashboard": {
    "title": "HashInsight Production Monitoring",
    "panels": [
      {
        "title": "Request Rate",
        "targets": [
          {
            "expr": "rate(hashinsight_requests_total[5m])"
          }
        ],
        "type": "timeseries"
      },
      {
        "title": "P95 Latency",
        "targets": [
          {
            "expr": "histogram_quantile(0.95, rate(hashinsight_request_latency_seconds_bucket[5m]))"
          }
        ],
        "type": "timeseries"
      },
      {
        "title": "Error Rate",
        "targets": [
          {
            "expr": "rate(hashinsight_requests_total{status=~\"5..\"}[5m]) / rate(hashinsight_requests_total[5m])"
          }
        ],
        "type": "timeseries"
      },
      {
        "title": "SLO Compliance",
        "targets": [
          {
            "expr": "sum(hashinsight_requests_total{status=200}) / sum(hashinsight_requests_total) * 100"
          }
        ],
        "type": "timeseries"
      }
    ]
  }
}
```

```

        "expr": "hashinsight_slo_compliance"
    }]
}
]
}
}

```

## 4.4 Alert Rules

### Prometheus Alert Rules

```

# prometheus/alerts.yml
groups:
- name: hashinsight_alerts
  interval: 30s
  rules:
    # Availability alert
    - alert: HighErrorRate
      expr: |
        rate(hashinsight_requests_total{status=~"5.."}[5m])
        / rate(hashinsight_requests_total[5m]) > 0.01
      for: 5m
      labels:
        severity: critical
      annotations:
        summary: "High 5xx error rate detected"
        description: "Error rate is {{ $value | humanizePercentage }}"

    # Latency alert
    - alert: HighLatency
      expr: |
        histogram_quantile(0.95,
          rate(hashinsight_request_latency_seconds_bucket[5m])
        ) > 0.25
      for: 10m
      labels:
        severity: warning
      annotations:
        summary: "P95 latency exceeds SL0"
        description: "P95 latency is {{ $value }}s (SL0: 0.25s)"

    # SL0 error budget alert
    - alert: ErrorBudgetExhausted
      expr: hashinsight_slo_error_budget_remaining < 0.1
      for: 5m

```

```

labels:
  severity: critical
annotations:
  summary: "SLO error budget nearly exhausted"
  description: "Only {{ $value | humanizePercentage }} budget remaining"

# Database connection alert
- alert: DatabaseConnectionFailure
  expr: hashinsight_db_connection_status == 0
  for: 2m
  labels:
    severity: critical
  annotations:
    summary: "Database connection failure"
    description: "Unable to connect to PostgreSQL"

# Cache hit rate alert
- alert: LowCacheHitRate
  expr: hashinsight_cache_hit_rate < 50
  for: 15m
  labels:
    severity: warning
  annotations:
    summary: "Low cache hit rate"
    description: "Cache hit rate is {{ $value }}%"

```

## 4.5 Circuit Breaker Configuration

HashInsight uses circuit breaker pattern to prevent cascading failures:

```

# monitoring/circuit_breaker.py
from monitoring.circuit_breaker import CircuitBreaker, circuit_breaker

# Database query circuit breaker
db_breaker = CircuitBreaker(
    failure_threshold=5,          # Trigger after 5 consecutive failures
    recovery_timeout=60,          # Attempt recovery after 60 seconds
    name="database_queries"
)

# API call circuit breaker
@circuit_breaker(
    failure_threshold=3,
    recovery_timeout=30,
    name="external_api"
)

```

```
def call_external_api():
    response = requests.get("https://api.coinwarz.com/.")
    return response.json()
```

## Circuit Breaker Status Monitoring

```
# View circuit breaker status
curl http://localhost:5000/api/circuit-breakers

# Response example
{
  "database_queries": {
    "state": "closed",
    "failure_count": 0,
    "total_calls": 1234,
    "success_rate": "99.8%"
  },
  "external_api": {
    "state": "half_open",
    "failure_count": 3,
    "total_calls": 456,
    "success_rate": "95.2%"
  }
}
```

## 4.6 Alert Notifications

### Slack Integration

```
# Environment variable configuration
SLACK_WEBHOOK_URL=https://hooks.slack.com/services/T00/B00/xxxxxx
ALERT_SLACK_CHANNEL=#hashinsight-alerts
```

### PagerDuty Integration

```
PAGERDUTY_API_KEY=xxxxxx
PAGERDUTY_SERVICE_KEY=xxxxxx
```

## 4.7 Monitoring Checklist

- [ ] Prometheus scraping metrics ([:9090/targets](#))
- [ ] Grafana dashboard displaying data
- [ ] Alert rules loaded
- [ ] Slack/PagerDuty notifications working
- [ ] SLO monitoring panel shows green
- [ ] Circuit breaker status normal

# Chapter 5: Backup & Recovery

## 5.1 Automated Backup Strategy

HashInsight uses [backup/backup\\_manager.py](#) for automated backups:

### Backup Features

- **PostgreSQL full backup** (pg\_dump)
- **AES-256 encryption** (backup file encryption)
- **gzip compression** (save storage space)
- **Remote storage** (S3/Azure/GCS support)
- **Integrity verification** (SHA256 checksum)

### Backup Scheduling

```
# Configure automatic backups via cron
# /etc/cron.d/hashinsight-backup

# Daily full backup at 2 AM
0 2 * * * /usr/bin/python3 /app/backup/backup_manager.py --type full

# Incremental backup every 4 hours
0 */4 * * * /usr/bin/python3 /app/backup/backup_manager.py --type incremental

# Upload to remote storage every Sunday at 3 AM
0 3 * * 0 /usr/bin/python3 /app/backup/backup_manager.py --upload
```

## 5.2 Manual Backup

```
# Execute full backup
python backup/backup_manager.py

# Backup output example
Backup created successfully: hashinsight_backup_20251003_140000.sql.gz.enc
Size: 245.3 MB
Encrypted: AES-256
Checksum: e3b0c44298fc1c149afbf4c8996fb92427ae41e4649b934ca495991b7852b855
Duration: 45.2s
```

## 5.3 Backup Retention Policy

Backup Type	Retention Period	Frequency
Full backup	30 days	Daily
Incremental backup	7 days	Every 4 hours
Weekly backup	12 weeks	Every Sunday
Monthly backup	12 months	1st of month

### Automatic Cleanup of Old Backups

```
# Clean up backups older than 30 days
python backup/backup_manager.py --cleanup --days 30

# Output
Deleting old backups: hashinsight_backup_20250903_*.sql.gz.enc
Cleanup complete: 2.1 GB freed
```

## 5.4 RTO/RPO Targets

Metric	Target	Actual
<b>RTO</b> (Recovery Time Objective)	≤4 hours	~2 hours
<b>RPO</b> (Recovery Point Objective)	≤15 minutes	~4 hours

## 5.5 Backup Recovery Process

### Step 1: List Available Backups

```
# List local backups
ls -lh /tmp/backups/

# List remote backups (S3)
aws s3 ls s3://hashinsight-backups/
```

### Step 2: Download Backup (if remote)

```
# Download from S3
aws s3 cp s3://hashinsight-backups/hashinsight_backup_20251003_020000.sql.gz.enc /tmp/restore/

# Download from Azure
az storage blob download \
    --account-name hashinsight \
    --container-name backups \
    --name hashinsight_backup_20251003_020000.sql.gz.enc \
    --file /tmp/restore/backup.sql.gz.enc
```

### Step 3: Decrypt Backup

```
# Decrypt using backup_manager
python backup/backup_manager.py \
    --decrypt /tmp/restore/hashinsight_backup_20251003_020000.sql.gz.enc \
    --output /tmp/restore/backup.sql.gz

# Manual decryption (if needed)
openssl enc -d -aes-256-cbc \
    -in hashinsight_backup_20251003_020000.sql.gz.enc \
```

```
-out backup.sql.gz \
-pass env:BACKUP_ENCRYPTION_KEY
```

## Step 4: Decompress Backup

```
# Decompress
gunzip /tmp/restore/backup.sql.gz

# Verify file integrity
sha256sum /tmp/restore/backup.sql
```

## Step 5: Restore Database

```
# ⚠ Warning: This will overwrite current database!

# 1. Stop application
systemctl stop hashinsight

# 2. Create emergency backup of current state
python backup/backup_manager.py --type emergency

# 3. Restore from backup
psql $DATABASE_URL < /tmp/restore/backup.sql

# 4. Verify restoration
psql $DATABASE_URL -c "SELECT COUNT(*) FROM miners;"

# 5. Restart application
systemctl start hashinsight
```

## 5.6 Disaster Recovery Plan

### Complete System Recovery Steps

```
# 1. Prepare new environment
export DATABASE_URL="new_database_url"
export SESSION_SECRET="new_session_secret"
export ENCRYPTION_PASSWORD="new_encryption_password"

# 2. Download latest backup
python backup/backup_manager.py --download-latest
```

```

# 3. Restore database
python backup/backup_manager.py --restore-latest

# 4. Deploy application
git clone https://github.com/hashinsight/platform.git
cd platform
pip install -r requirements.txt

# 5. Verify configuration
python scripts/verify_setup.py

# 6. Start application
gunicorn --bind 0.0.0.0:5000 --workers 4 main:app

# 7. Verify service health
curl http://localhost:5000/health

```

## 5.7 Backup Testing

```

# Regular backup restoration testing (quarterly)
# 1. Create test environment
docker run -d --name test-postgres -e POSTGRES_PASSWORD=test postgres:15

# 2. Restore to test environment
export TEST_DATABASE_URL="postgresql://postgres:test@localhost:5432/test_db"
psql $TEST_DATABASE_URL < /tmp/backups/latest_backup.sql

# 3. Verify data integrity
python scripts/verify_backup_integrity.py --database $TEST_DATABASE_URL

# 4. Cleanup test environment
docker stop test-postgres
docker rm test-postgres

```

## 5.8 Backup Checklist

- [ ] Daily automated backups running
- [ ] Backup encryption enabled
- [ ] Remote storage configured
- [ ] Backup integrity verification passing
- [ ] Restoration tested within last 90 days
- [ ] Old backups automatically cleaned up

- [ ] Backup monitoring alerts configured
  - [ ] Documentation up to date
- 

## Chapter 6: Security Operations

---

### 6.1 KMS Key Management

HashInsight supports enterprise-grade Key Management Systems (KMS):

#### Supported KMS Providers

- **AWS KMS:** Amazon Web Services Key Management Service
- **GCP KMS:** Google Cloud Platform Key Management Service
- **Azure Key Vault:** Microsoft Azure Key Vault

#### AWS KMS Configuration

```
# Configure AWS KMS
from security.kms_client import KMSClient, KMSProvider, EncryptionContext

kms_config = {
    'key_id': 'arn:aws:kms:us-east-1:123456789:key/xxxxx',
    'region': 'us-east-1'
}

client = KMSClient(KMSProvider.AWS_KMS, kms_config)

# Encrypt data
context = EncryptionContext(
    purpose="user_data_encryption",
    tenant_id="tenant_123"
)

ciphertext = client.encrypt_secret(
    plaintext="sensitive data",
    key_id=kms_config['key_id'],
    context=context
)
```

## GCP KMS Configuration

```
# Configure GCP KMS
kms_config = {
    'project_id': 'hashinsight-prod',
    'location': 'us-east1',
    'keyring': 'hashinsight-keyring',
    'key_id': 'encryption-key'
}

client = KMSClient(KMSProvider.GCP_KMS, kms_config)
```

## Azure Key Vault Configuration

```
# Configure Azure Key Vault
kms_config = {
    'vault_url': 'https://hashinsight-vault.vault.azure.net/',
    'key_name': 'encryption-key'
}

client = KMSClient(KMSProvider.AZURE_KEY_VAULT, kms_config)
```

## Key Rotation Process

```
# 1. Create new key version in KMS
aws kms create-key --description "HashInsight Master Key v2"

# 2. Update application configuration
export AWS_KMS_KEY_ID=arn:aws:kms:us-east-1:123456789:key/new-key-id

# 3. Re-encrypt existing data (background task)
python scripts/rotate_encryption_keys.py --old-key OLD_KEY_ID --new-key NEW_KEY_ID

# 4. Verify new key
python scripts/verify_encryption.py --key-id NEW_KEY_ID

# 5. Disable old key (retain for 90 days)
aws kms disable-key --key-id OLD_KEY_ID
```

## 6.2 mTLS Mutual Authentication

### Certificate Generation

```
# 1. Generate CA certificate
openssl genrsa -out ca.key 4096
openssl req -new -x509 -days 3650 -key ca.key -out ca.crt \
-subj "/C=US/O=HashInsight/CN=HashInsight Root CA"

# 2. Generate server certificate
openssl genrsa -out server.key 4096
openssl req -new -key server.key -out server.csr \
-subj "/C=US/O=HashInsight/CN=*.hashinsight.net"
openssl x509 -req -days 365 -in server.csr -CA ca.crt -CAkey ca.key \
-CAcreateserial -out server.crt

# 3. Generate client certificate
openssl genrsa -out client.key 4096
openssl req -new -key client.key -out client.csr \
-subj "/C=US/O=HashInsight/CN=client.hashinsight.net"
openssl x509 -req -days 365 -in client.csr -CA ca.crt -CAkey ca.key \
-CAcreateserial -out client.crt
```

### mTLS Configuration

```
# Environment variable configuration
export MTLS_ENABLED=true
export MTLS_CA_CERT_PATH=/app/certs/ca.crt
export MTLS_SERVER_CERT_PATH=/app/certs/server.crt
export MTLS_SERVER_KEY_PATH=/app/certs/server.key
export MTLS_VERIFY_CLIENT=true
export MTLS_ALLOWED_DN_PATTERNS="CN=*.hashinsight.net,O=HashInsight,C=US"
```

### Using mTLS Authentication

```
from common.mtls_auth import require_mtls

@app.route('/api/admin/sensitive')
@require_mtls()
def sensitive_endpoint():
    # Only allow requests with valid client certificates
```

```
client_dn = g.client_cert_subject  
return jsonify({"message": f"Authenticated as {client_dn}"})
```

## 6.3 API Key Management

### Key Format

HashInsight API key format: `hsı_{env}_key_{random}`

Examples: - Production: `hsı_prod_key_a1b2c3d4e5f6g7h8` - Development:

`hsı_dev_key_x9y8z7w6v5u4t3s2`

### Create API Key

```
# Create using management tool  
python scripts/create_api_key.py \  
--user-id 123 \  
--permissions "miners:read,miners:write" \  
--expires-in 90  
  
# Output  
API key created  
Key: hsi_prod_key_a1b2c3d4e5f6g7h8  
User ID: 123  
Permissions: miners:read,miners:write  
Expires: 2025-12-31T23:59:59Z  
⚠ Please store this key securely, it will not be shown again
```

### API Key Rotation

```
# 1. Create new key  
new_key=$(python scripts/create_api_key.py --user-id 123 --copy-from old_key_id)  
  
# 2. Update client configuration (7-day dual-key period)  
# Both old and new keys valid simultaneously  
  
# 3. Verify new key  
curl -H "Authorization: Bearer $new_key" http://localhost:5000/api/miners  
  
# 4. Revoke old key  
python scripts/revoke_api_key.py --key-id old_key_id
```

## 6.4 WireGuard Enterprise Network

### Hub Server Deployment

```
# 1. Run Hub installation script
sudo bash wireguard/hub_setup.sh

# 2. Configure firewall
sudo ufw allow 51820/udp
sudo ufw enable

# 3. Start WireGuard
sudo systemctl enable wg-quick@wg0
sudo systemctl start wg-quick@wg0

# 4. Check status
sudo wg show
```

### Site Gateway Configuration

```
# 1. Generate site keys
cd wireguard/site-gateway
python key_manager.py --generate-keys --site beijing-dc1

# 2. Add to Hub configuration
sudo nano /etc/wireguard/wg0.conf

# Add Peer configuration
[Peer]
PublicKey = site_public_key
AllowedIPs = 10.8.1.0/24
Endpoint = beijing-gateway.hashinsight.net:51820

# 3. Reload configuration
sudo wg-quick down wg0
sudo wg-quick up wg0
```

## 6.5 Audit Logging

HashInsight logs all critical operations to audit logs ([audit/audit\\_logger.py](#)):

## Audit Event Types

- Authentication (login/logout/failure)
- Data access (CRUD operations)
- Configuration changes
- Permission changes
- Encryption operations
- API key management
- Suspicious activities

## Audit Log Format

```
{  
  "timestamp": "2025-10-03T12:00:00.000Z",  
  "event_id": "a1b2c3d4e5f6g7h8",  
  "level": "INFO",  
  "category": "authentication",  
  "action": "login",  
  "user_id": "123",  
  "user_email": "user@example.com",  
  "user_role": "admin",  
  "ip_address": "192.168.1.100",  
  "status": "success",  
  "details": {  
    "login_method": "password",  
    "two_factor": true  
  }  
}
```

## Query Audit Logs

```
# View last 100 audit logs  
tail -100 logs/audit.jsonl  
  
# Query specific user operations  
jq 'select(.user_email == "admin@hashinsight.net")' logs/audit.jsonl  
  
# Query failed login attempts  
jq 'select(.action == "login_failed")' logs/audit.jsonl  
  
# Query security events from last 24 hours  
jq --arg date "$(date -d '24 hours ago' -Iseconds)" \  
  'select(.timestamp >=$date)' logs/audit.jsonl
```

```
'select(.timestamp > $date and .level == "SECURITY")' \
logs/audit.jsonl
```

## 6.6 Compliance Requirements

HashInsight follows these compliance standards:

### SOC 2 Type II

- Access control
- Encrypted transmission (TLS 1.3)
- Encrypted storage (AES-256)
- Audit logging (immutable)
- Change management
- Disaster recovery

### PCI DSS (if processing payments)

- Sensitive data masking
- Key management (KMS)
- Network isolation (WireGuard)
- Regular penetration testing

### GDPR

- Data minimization
- User data export
- Data deletion (right to be forgotten)
- Data processing records

## 6.7 Security Checklist

- [ ] All keys stored in KMS (not in code/config files)
- [ ] mTLS enabled (production environment)
- [ ] API keys rotated regularly (90 days)
- [ ] SSL/TLS certificates valid and not expired
- [ ] Audit logs writing normally
- [ ] No sensitive information leaked to logs

- [ ] Database connections encrypted (SSL)
  - [ ] Session key strength  $\geq$ 256 bits
  - [ ] Regular security scans (quarterly)
  - [ ] Penetration testing report (annually)
- 

## Chapter 7: Troubleshooting Guide

---

### 7.1 Application Won't Start

#### Symptoms

```
$ gunicorn main:app
[ERROR] Application failed to start
```

#### Diagnostic Steps

##### 1. Check Port Usage

```
# Check port 5000 usage
lsof -i :5000

# If occupied, kill process
kill -9 $(lsof -t -i:5000)
```

##### 2. Check Environment Variables

```
# Verify required variables
python3 << 'EOF'
import os
required = ['DATABASE_URL', 'SESSION_SECRET', 'ENCRYPTION_PASSWORD']
for var in required:
    value = os.getenv(var)
    if not value:
        print(f"  Missing: {var}")
    else:
        print(f"  {var}: {'*' * 8} (set)")
EOF
```

### **3. Check Database Connection**

```
# Test PostgreSQL connection
psql $DATABASE_URL -c "SELECT version();"

# If fails, check Neon endpoint status
# Visit https://console.neon.tech
```

### **4. Check Python Dependencies**

```
# Verify key libraries
python3 -c "
import flask
import gunicorn
import psycopg2
import sqlalchemy
print(' All dependencies OK')
"
```

### **5. View Detailed Errors**

```
# Enable debug mode
export FLASK_DEBUG=1
python3 main.py

# View full stack trace
```

## Common Errors and Solutions

Error	Cause	Solution
<code>ModuleNotFoundError: No module named 'flask'</code>	Dependencies not installed	<code>pip install -r requirements.txt</code>
<code>OperationalError: could not connect to server</code>	Database connection failed	Check DATABASE_URL, verify Neon endpoint
<code>ValueError: SECRET_KEY must be set</code>	SESSION_SECRET not set	<code>export SESSION_SECRET=\$(python -c 'import secrets; print(secrets.token_urlsafe(32))')</code>
<code>Address already in use</code>	Port occupied	<code>kill -9 \$(lsof -t -i:5000)</code>

## 7.2 Database Connection Failure

### Symptoms

```
sqlalchemy.exc.OperationalError: (psycopg2.OperationalError)
could not connect to server: Connection timed out
```

### Diagnostic Steps

#### 1. Verify Connection String

```
# Parse DATABASE_URL
python3 << EOF
import os
from urllib.parse import urlparse
url = os.getenv('DATABASE_URL')
parsed = urlparse(url)
print(f"Host: {parsed.hostname}")
print(f"Port: {parsed.port}")
print(f"Database: {parsed.path[1:]}")
print(f"User: {parsed.username}")
EOF
```

#### 2. Test Network Connectivity

```

# Extract host and port
export DB_HOST=$(python3 -c "from urllib.parse import urlparse; import os; print(urlparse(os.getenv('DB_URL')).hostname)")
export DB_PORT=$(python3 -c "from urllib.parse import urlparse; import os; print(urlparse(os.getenv('DB_URL')).port)")

# Test TCP connection
nc -zv $DB_HOST $DB_PORT

```

### 3. Check Connection Pool

```

# Connection pool configuration (config.py)
SQLALCHEMY_ENGINE_OPTIONS = {
    'pool_size': 10,           # Default 10 connections
    'pool_recycle': 300,       # Recycle after 5 minutes
    'pool_pre_ping': True,    # Test connection before use
    'pool_timeout': 30,        # 30 second timeout
    'max_overflow': 20,        # Max 20 overflow connections
    'connect_args': {
        'connect_timeout': 15 # 15 second connect timeout
    }
}

```

### 4. Check Neon Endpoint Status

```

# Visit Neon Console
# https://console.neon.tech/app/projects

# Check endpoint status:
# - Active (green) - Normal
# - Idle (yellow) - Needs wake-up
# - Suspended (gray) - Suspended

```

### 5. Enable Connection Retry

```

# Add retry logic in app.py
from sqlalchemy import event, exc
from sqlalchemy.pool import Pool

@event.listens_for(Pool, "connect")
def receive_connect(dbapi_conn, connection_record):
    connection_record.info['pid'] = os.getpid()

@event.listens_for(Pool, "checkout")
def receive_checkout(dbapi_conn, connection_record, connection_proxy):

```

```
pid = os.getpid()
if connection_record.info['pid'] != pid:
    connection_record.connection = connection_proxy.connection = None
    raise exc.DisconnectionError(
        "Connection record belongs to pid %s, "
        "attempting to check out in pid %s" %
        (connection_record.info['pid'], pid)
)
```

## 7.3 Cache Issues

### Symptoms

- Response time significantly increased
- Cache hit rate below 50%
- Redis connection errors

### Diagnostic Steps

#### 1. Check Redis Connection

```
# Test Redis
redis-cli ping
# Expected: PONG

# Check Redis memory
redis-cli info memory

# Check key count
redis-cli dbsize
```

#### 2. View Cache Hit Rate

```
# Query via API
curl http://localhost:5000/api/cache/stats

# Expected response
{
    "cache_type": "redis",
    "hit_rate": 75.5,
    "total_requests": 10000,
    "hits": 7550,
```

```
        "misses": 2450
    }
```

### 3. Cache Fallback Mechanism

```
# cache_manager.py automatic fallback
if redis_available:
    cache_backend = RedisCache()
else:
    logger.warning("Redis unavailable, falling back to memory cache")
    cache_backend = MemoryCache()
```

### 4. Clear Cache

```
# Clear all cache
redis-cli FLUSHDB

# Clear specific prefix
redis-cli --scan --pattern 'hashinsight:' | xargs redis-cli DEL
```

## 7.4 Performance Degradation

### Symptoms

- API response time p95 > 250ms
- Slow database queries
- High CPU/memory usage

### Diagnostic Steps

#### 1. Check Request Coalescing Status

```
# View request merging statistics
curl http://localhost:5000/api/performance/coalescing-stats

# Expected response
{
    "enabled": true,
    "performance_improvement": "9.8x",
    "deduplicated_requests": 5432,
```

```
        "total_requests": 53210
    }
```

## 2. Analyze Slow Queries

```
-- Enable slow query log (PostgreSQL)
ALTER DATABASE hashinsight_db SET log_min_duration_statement = 1000;

-- Query slow queries
SELECT query, calls, total_time, mean_time
FROM pg_stat_statements
WHERE mean_time > 1000
ORDER BY mean_time DESC
LIMIT 10;
```

## 3. Check Database Indexes

```
# Run index optimization
python database/optimize_indexes.py --analyze

# View missing indexes
python database/optimize_indexes.py --suggest
```

## 4. Monitor System Resources

```
# CPU usage
top -b -n 1 | grep gunicorn

# Memory usage
ps aux | grep gunicorn | awk '{sum+=$6} END {print sum/1024 " MB"}'

# Database connections
psql $DATABASE_URL -c "SELECT count(*) FROM pg_stat_activity WHERE datname = 'hashinsight_db';"
```

## 5. Enable Performance Profiling

```
# Add to routes requiring profiling
from werkzeug.middleware.profiler import ProfilerMiddleware

app.wsgi_app = ProfilerMiddleware(
    app.wsgi_app,
    restrictions=[10],
```

```
    profile_dir='./profiles'  
)
```

## 7.5 Blockchain Integration Errors

### Symptoms

```
ERROR:blockchain_integration: Encryption configuration error: ENCRYPTION_PASSWORD environment variable  
ERROR:sla_nft_routes: Failed to get SLA status
```

### Diagnostic Steps

#### 1. Check ENCRYPTION\_PASSWORD

```
# Verify environment variable  
echo $ENCRYPTION_PASSWORD  
  
# If not set  
export ENCRYPTION_PASSWORD=$(python3 -c 'import secrets; print(secrets.token_urlsafe(32))')
```

#### 2. Check Blockchain Configuration

```
# Verify blockchain environment variables  
python3 << 'EOF'  
import os  
config = {  
    'BLOCKCHAIN_ENABLED': os.getenv('BLOCKCHAIN_ENABLED'),  
    'BLOCKCHAIN_PRIVATE_KEY': '****' if os.getenv('BLOCKCHAIN_PRIVATE_KEY') else None,  
    'BLOCKCHAIN_NETWORK': os.getenv('BLOCKCHAIN_NETWORK', 'base-sepolia'),  
    'BASE_RPC_URL': os.getenv('BASE_RPC_URL', 'https://sepolia.base.org')  
}  
for k, v in config.items():  
    status = ' ' if v else ''  
    print(f"{status} {k}: {v}")  
EOF
```

#### 3. Test Web3 Connection

```
python3 << 'EOF'  
from web3 import Web3  
import os
```

```

rpc_url = os.getenv('BASE_RPC_URL', 'https://sepolia.base.org')
w3 = Web3(HTTPProvider(rpc_url))

if w3.is_connected():
    print(f"  Web3 connected to {rpc_url}")
    print(f"Block number: {w3.eth.block_number}")
else:
    print(f"  Web3 connection failed")
EOF

```

#### 4. Verify Private Key Format

```

# Private key should start with 0x, followed by 64 hex characters
python3 << 'EOF'
import os
import re

private_key = os.getenv('BLOCKCHAIN_PRIVATE_KEY', '')
if re.match(r'^0x[0-9a-fA-F]{64}$', private_key):
    print("  Private key format valid")
else:
    print("  Invalid private key format")
    print("Expected: 0x + 64 hex characters")
EOF

```

## 7.6 Log Locations

Log Type	Path	Format
Application logs	stdout	Text
Audit logs	logs/audit.jsonl	JSON Lines
Error logs	stderr	Text
Gunicorn logs	/var/log/gunicorn/	Text
PostgreSQL logs	Neon Console	Text
Workflow logs	/tmp/logs/	Text

## View Logs

```
# Real-time application logs
tail -f /var/log/hashinsight/app.log

# View error logs
grep ERROR /var/log/hashinsight/app.log | tail -50

# View audit logs
tail -f logs/audit.jsonl | jq '.'

# View Gunicorn access logs
tail -f /var/log/gunicorn/access.log
```

# Chapter 8: Daily Operations

## 8.1 Database Maintenance

### 8.1.1 Index Optimization

```
# Automatically optimize indexes
python database/optimize_indexes.py --auto

# Analyze and suggest indexes
python database/optimize_indexes.py --analyze --suggest

# Output example:
#   Existing indexes: 45
#   Scanning slow queries...
#   Suggested indexes:
#     - CREATE INDEX idx_miners_user_id ON miners(user_id)
#     - CREATE INDEX idx_market_data_timestamp ON market_analytics(created_at DESC)
```

### 8.1.2 Connection Pool Monitoring

```
-- View current connections
SELECT
    pid,
    username,
```

```

application_name,
client_addr,
state,
query_start
FROM pg_stat_activity
WHERE datname = 'hashinsight_db'
ORDER BY query_start;

-- Kill idle connections
SELECT pg_terminate_backend(pid)
FROM pg_stat_activity
WHERE datname = 'hashinsight_db'
AND state = 'idle'
AND state_change < NOW() - INTERVAL '30 minutes';

```

### 8.1.3 Slow Query Analysis

```

-- Enable pg_stat_statements
CREATE EXTENSION IF NOT EXISTS pg_stat_statements;

-- View Top 10 slow queries
SELECT
    substring(query, 1, 100) AS short_query,
    calls,
    total_time,
    mean_time,
    max_time
FROM pg_stat_statements
WHERE query NOT LIKE '%pg_stat_statements%'
ORDER BY mean_time DESC
LIMIT 10;

-- Reset statistics
SELECT pg_stat_statements_reset();

```

### 8.1.4 Database Cleanup

```

# Clean old data (retain 90 days)
python scripts/cleanup_old_data.py --days 90

# Cleanup example:
#   Deleted market_analytics older than 90 days: 12,543 records

```

```
# Deleted audit logs older than 90 days: 45,123 records
# Cleanup complete, freed space: 2.3 GB
```

## 8.2 Cache Management

### 8.2.1 Redis Cache Cleanup

```
# Clear all HashInsight cache
redis-cli --scan --pattern 'hashinsight:*' | xargs redis-cli DEL

# Clear specific module cache
redis-cli --scan --pattern 'hashinsight:miners:' | xargs redis-cli DEL

# Clear expired keys (Redis automatic, manual trigger)
redis-cli --scan --pattern 'hashinsight:*' | while read key; do
    redis-cli TTL "$key"
done
```

### 8.2.2 Memory Cache Monitoring

```
# Query memory cache status via API
curl http://localhost:5000/api/cache/memory-stats

# Response
{
  "cache_type": "memory",
  "size_mb": 124.5,
  "entries": 5432,
  "hit_rate": 82.3,
  "evictions": 234
}
```

### 8.2.3 Request Coalescer Status

```
# View request merging statistics
curl http://localhost:5000/api/performance/coalescing-stats | jq

# Manually clear coalescing cache
curl -X POST http://localhost:5000/api/performance/clear-coalescing-cache
```

## 8.3 Batch Tasks

### 8.3.1 Bulk Import Miners

```
# Prepare CSV file (max 5000 miners)
# Format: name,model,hashrate_th,power_w,efficiency

# Execute bulk import
python batch/batch_import_manager.py \
--file miners_upload_5000.csv \
--user-id 123 \
--validate

# Output:
#   File validation passed
#   Record count: 5000
#   Processing...
# [██████████] 100%
#   Import complete: 5000 miners (Duration: 45.2s)
```

### 8.3.2 Data Collection Tasks

HashInsight automatically collects market data every 15 minutes:

```
# Manually trigger data collection
python modules/analytics/engines/analytics_engine.py --collect-now

# View collection status
curl http://localhost:5000/api/analytics/collection-status

# Response
{
  "last_collection": "2025-10-03T12:15:00Z",
  "next_collection": "2025-10-03T12:30:00Z",
  "status": "healthy",
  "data_points_today": 8
}
```

### 8.3.3 Data Collection Scheduling

```
# Configure data collection frequency (config.py)
ANALYTICS_COLLECTION_INTERVAL = 15 # minutes
ANALYTICS_MAX_DATA_POINTS_PER_DAY = 10 # Daily limit
```

```
# Enable/disable background services (default: disabled, i.e., 0)
export ENABLE_BACKGROUND_SERVICES=1 # Enable background data collection
export ENABLE_BACKGROUND_SERVICES=0 # Disable background data collection (default)
```

## 8.4 User Management

### 8.4.1 Create User

```
# Create via admin interface
# Visit: http://localhost:5000/admin/users/create

# Create via command line
python scripts/create_user.py \
--email admin@hashinsight.net \
--username admin \
--role owner \
--password-prompt

# Output:
# Please enter password: *****
# User created
# ID: 123
# Email: admin@hashinsight.net
# Role: owner
```

### 8.4.2 Role Assignment

Role	Permissions	Description
owner	All permissions	System owner
admin	Management permissions	System administrator
broker	Broker permissions	Customer management, orders
client	Client permissions	View own data

```
# Change user role
python scripts/change_user_role.py \
--user-id 123 \
```

```
--new-role admin

# Bulk import users
python scripts/bulk_import_users.py --file users.csv
```

### 8.4.3 Permission Management

```
# Check user permissions
from decorators import has_permission

@app.route('/api/miners/delete/<int:miner_id>', methods=['DELETE'])
@login_required
@requires_permission('miners:delete')
def delete_miner(miner_id):
    # Only allow users with miners:delete permission
    pass
```

## 8.5 Data Cleanup

### 8.5.1 Daily Data Point Limits

To control storage costs, HashInsight limits daily data points:

```
# config.py
ANALYTICS_MAX_DATA_POINTS_PER_DAY = 10 # Max 10 data points per day
```

### 8.5.2 Historical Data Archival

```
# Archive data older than 90 days to cold storage
python scripts/archive_historical_data.py --days 90 --storage s3

# Output:
#   Scanning data for archival...
#   Archived 45,234 records to S3
#   Deleted archived data from primary database
#   Duration: 12m 34s
```

## 8.6 Monitoring Tasks

### 8.6.1 Daily Health Check

```
# Run daily health check script
python scripts/daily_health_check.py

# Output:
#   Application: Healthy
#   Database: Connected, 15 active connections
#   Cache: Redis available, 78.5% hit rate
#   Disk space: 234 GB available (85% free)
#   Backups: Latest backup 2 hours ago
#   SSL certificates: Valid, expires in 45 days
# △ Warning: API key 'hs1_prod_key_abc123' expires in 7 days
```

### 8.6.2 Weekly Reports

```
# Generate weekly operations report
python scripts/weekly_report.py --email ops@hashinsight.net

# Report includes:
# - System uptime and availability
# - Performance metrics (latency, error rate)
# - Resource usage trends
# - Security events
# - Backup status
# - Action items
```

## Chapter 9: Performance Tuning Guide

### 9.1 Request Coalescing Optimization

HashInsight implements Request Coalescing pattern to achieve **9.8x performance improvement**:

## How It Works

```
# Multiple concurrent identical requests merged into single request
# Example: 10 simultaneous "get BTC price" requests → 1 API call

from performance.request_coalescer import RequestCoalescer

coalescer = RequestCoalescer()

@coalescer.coalesce(key='btc_price')
def get_btc_price():
    # Expensive API call
    response = requests.get('https://api.coingecko.com/...')
    return response.json()

# 10 concurrent calls to get_btc_price() → Only 1 API call executed
# 9 requests wait for the first one to complete and share result
```

## Performance Metrics

```
# View coalescing statistics
curl http://localhost:5000/api/performance/coalescing-stats

# Response
{
  "enabled": true,
  "total_requests": 98234,
  "deduplicated_requests": 85432,
  "api_calls_saved": 85432,
  "performance_improvement": "9.8x",
  "average_wait_time_ms": 45
}
```

## 9.2 Database Optimization

### 9.2.1 Index Strategy

```
-- High-frequency query field indexes
CREATE INDEX idx_miners_user_id ON miners(user_id);
CREATE INDEX idx_market_data_timestamp ON market_analytics(created_at DESC);
CREATE INDEX idx_calculations_user_created ON calculations(user_id, created_at);
```

```
-- Composite indexes (commonly queried together)
CREATE INDEX idx_miners_user_model ON miners(user_id, model);

-- Partial indexes (index only active data)
CREATE INDEX idx_active_miners ON miners(user_id) WHERE status = 'active';
```

## 9.2.2 Query Optimization

### Use EXPLAIN to Analyze Queries

```
EXPLAIN ANALYZE
SELECT m.*, u.username
FROM miners m
JOIN users u ON m.user_id = u.id
WHERE m.status = 'active'
ORDER BY m.created_at DESC
LIMIT 100;

-- Review execution plan, optimize slow queries
```

### Avoid N+1 Queries

```
# Bad - N+1 queries
miners = Miner.query.filter_by(user_id=user_id).all()
for miner in miners:
    print(miner.user.username) # One query per loop iteration

# Good - Use JOIN
miners = Miner.query.join(User).filter(Miner.user_id == user_id).all()
for miner in miners:
    print(miner.user.username) # Single query
```

## 9.2.3 Connection Pool Tuning

```
# config.py
SQLALCHEMY_ENGINE_OPTIONS = {
    'pool_size': 10,           # Adjust based on concurrency (workers × 2)
    'pool_recycle': 300,       # Recycle connections after 5 minutes
    'pool_pre_ping': True,     # Test connection before use (Neon required)
    'pool_timeout': 30,        # 30 second timeout
    'max_overflow': 20,        # Max 20 overflow connections
    'connect_args': {
        'connect_timeout': 15,   # Connection timeout
```

```

        'application_name': 'hashinsight', # For monitoring
        'options': '-c statement_timeout=30000' # Query timeout 30 seconds
    }
}

```

## 9.3 Caching Strategy

### 9.3.1 Multi-tier Caching

```

# Three-tier cache architecture
L1: Memory cache (fastest, smallest capacity)
    ↓ miss
L2: Redis cache (fast, medium capacity)
    ↓ miss
L3: Database (slowest, largest capacity)

# Implementation
@cache_manager.cached(ttl=300, level='L1') # 5 minutes
def get_user_profile(user_id):
    return db.session.query(User).get(user_id)

@cache_manager.cached(ttl=3600, level='L2') # 1 hour
def get_market_data():
    return fetch_from_api()

```

### 9.3.2 TTL Configuration

Data Type	TTL	Reason
User info	5 minutes	May change
Market data	5 minutes	Real-time requirement
Miner list	1 hour	Low change frequency
Statistics	10 minutes	Compute-intensive
Static content	24 hours	Rarely changes

### 9.3.3 Cache Warming

```
# Warm critical caches at startup
python scripts/warmup_cache.py

# Script contents:
# 1. Load popular user data
# 2. Load current market data
# 3. Pre-calculate statistics
# 4. Load configuration data

# Output:
#   Warmed user cache: 1,234 entries
#   Warmed market data: Current price + hashrate
#   Warmed statistics: 10 dashboards
#   Duration: 12.3s
```

## 9.4 Batch Processing Optimization

### 9.4.1 Vectorized Computation

```
# Bad - Individual calculation
for miner in miners:
    daily_revenue = calculate_revenue(
        miner.hashrate,
        btc_price,
        network_difficulty
    )
# Duration: 5000 miners × 10ms = 50 seconds

# Good - Vectorized calculation
import numpy as np

hashrates = np.array([m.hashrate for m in miners])
revenues = calculate_revenue_vectorized(
    hashrates,
    btc_price,
    network_difficulty
)
# Duration: ~500ms (100x improvement)
```

#### 9.4.2 Concurrent Processing

```
from concurrent.futures import ThreadPoolExecutor

# Concurrent batch processing
def process_miner_batch(miners, batch_size=100):
    with ThreadPoolExecutor(max_workers=8) as executor:
        futures = []
        for i in range(0, len(miners), batch_size):
            batch = miners[i:i+batch_size]
            future = executor.submit(process_batch, batch)
            futures.append(future)

    results = [f.result() for f in futures]
    return results
```

#### 9.4.3 Memory Optimization

```
# Batch insert (avoid ORM overhead)
from sqlalchemy import insert

# Bad - Individual inserts
for data in large_dataset:
    db.session.add(Model(**data))
    db.session.commit() # Commit each time

# Good - Batch insert
db.session.bulk_insert_mappings(Model, large_dataset)
db.session.commit() # Single commit

# Better - Use bulk insert
stmt = insert(Model).values(large_dataset)
db.session.execute(stmt)
db.session.commit()
```

### 9.5 Performance Monitoring

```
# Regularly run performance benchmarks
python scripts/performance_benchmark.py

# Output:
#   API response times:
#     - /api/miners: p50=45ms, p95=120ms, p99=250ms
```

```
#   - /api/dashboard: p50=80ms, p95=200ms, p99=400ms
# Database queries:
#   - Average query time: 15ms
#   - Slow queries (>1s): 0
# Cache performance:
#   - Hit rate: 78.5%
#   - Request Coalescing: 9.8x improvement
```

## Chapter 10: Emergency Response

---

### 10.1 On-Call Rotation

#### Rotation Schedule

Time Period	Primary On-Call	Backup On-Call	Escalation Contact
Weekdays 09:00-18:00	DevOps Engineer	Backend Engineer	Technical Director
Weekdays 18:00-09:00	Rotation Engineer	Backup Engineer	On-Call Manager
Weekends/Holidays	Rotation Engineer	Backup Engineer	On-Call Manager

#### On-Call Tools

- **PagerDuty:** Alert notifications
- **Slack:** #incident-response channel
- **Grafana:** Real-time monitoring
- **Incident.io:** Incident management

## 10.2 Incident Severity Levels

Level	Impact	Response Time	Escalation Time	Examples
P0	Complete service outage	15 minutes	30 minutes	Database down, application inaccessible
P1	Core functionality degraded	30 minutes	1 hour	API latency >5s, error rate >5%
P2	Partial functionality failure	2 hours	4 hours	Single feature unavailable
P3	Performance degradation	4 hours	8 hours	Slower response but functional
P4	Non-urgent issues	1 business day	2 business days	UI display issues, minor bugs

## 10.3 Response Process

### P0/P1 Critical Incident Process

1. [0-5 minutes] Alert triggered
  - ↳ PagerDuty notifies primary on-call
  - ↳ Automatically create Slack incident channel
  - ↳ Automatically notify backup on-call
2. [5-15 minutes] Initial response
  - ↳ Acknowledge incident (ACK alert)
  - ↳ Post initial status update
  - ↳ Begin troubleshooting
  - ↳ Record timeline
3. [15-30 minutes] Emergency mitigation
  - ↳ Implement temporary mitigation
  - ↳ Assess impact scope
  - ↳ Decide whether to escalate
  - ↳ Update incident status
4. [30 minutes+] Full resolution
  - ↳ Implement permanent fix

- ↳ Verify service recovery
- ↳ Post recovery announcement
- ↳ Begin post-incident review

## 10.4 Common Emergency Scenarios

### Scenario 1: Database Connection Pool Exhausted

#### Symptoms

```
sqlalchemy.exc.TimeoutError: QueuePool limit of size 10 overflow 20 reached
```

#### Emergency Response

```
# 1. Immediately kill long idle connections
psql $DATABASE_URL -c "
SELECT pg_terminate_backend(pid)
FROM pg_stat_activity
WHERE state = 'idle'
    AND state_change < NOW() - INTERVAL '10 minutes';
"

# 2. Increase connection pool limits (temporary)
export SQLALCHEMY_POOL_SIZE=20
export SQLALCHEMY_MAX_OVERFLOW=40

# 3. Restart application
systemctl restart hashinsight

# 4. Monitor connection count
watch -n 5 "psql $DATABASE_URL -c 'SELECT count(*) FROM pg_stat_activity;'"
```

### Scenario 2: Memory Leak Causing OOM

#### Symptoms

```
MemoryError: Unable to allocate array
Killed (OOM)
```

#### Emergency Response

```
# 1. Immediately restart affected worker  
kill -HUP $(cat /var/run/gunicorn.pid)  
  
# 2. Clear cache to free memory  
redis-cli FLUSHDB  
  
# 3. Limit worker count (temporary)  
gunicorn --workers 2 --max-requests 500 main:app  
  
# 4. Monitor memory usage  
watch -n 5 'free -h'
```

### Scenario 3: Malicious Traffic Attack

**Symptoms** - Abnormally high request volume - Many 401/403 errors - Suspicious requests from specific IP ranges

#### Emergency Response

```
# 1. Enable rate limiting  
export RATE_LIMIT_ENABLED=true  
export RATE_LIMIT_REQUESTS_PER_MINUTE=10  
  
# 2. Block malicious IPs (via firewall)  
ufw deny from 192.168.1.0/24  
  
# 3. Enable Cloudflare protection (if using)  
# Visit Cloudflare Dashboard -> Security -> DDoS  
  
# 4. Analyze attack patterns  
tail -1000 /var/log/gunicorn/access.log | \  
awk '{print $1}' | sort | uniq -c | sort -rn | head -20
```

## 10.5 Rollback Procedures

### Code Rollback

```
# 1. Determine rollback version  
git log --oneline -10  
  
# 2. Rollback to last stable version  
git revert HEAD --no-edit  
# Or
```

```
git checkout v1.9.5

# 3. Redeploy
git push origin main

# 4. Verify rollback success
curl http://localhost:5000/health | jq '.version'

# 5. Notify team
# Slack: "Rolled back to v1.9.5, service restored"
```

## Database Rollback

```
# △ Warning: Database rollback is high-risk, must backup first!

# 1. Create current state backup
python backup/backup_manager.py --type emergency

# 2. Restore to rollback point
psql $DATABASE_URL < /tmp/backups/hashinsight_backup_20251003_020000.sql

# 3. Verify data integrity
python scripts/verify_database_integrity.py

# 4. Restart application
systemctl restart hashinsight
```

## 10.6 Communication Templates

### Initial Incident Announcement

[P0 Incident] HashInsight Service Outage

Time: 2025-10-03 14:23 UTC

Impact: All users unable to access main application

Status: Under investigation

We have confirmed a service outage, team is urgently investigating.  
Will provide update within 30 minutes.

Incident channel: #incident-2025-10-03-db-outage

On-call engineer: @john.doe

## Progress Update

[Update] HashInsight Incident Progress

Time: 2025-10-03 14:45 UTC

Root cause: Database connection pool exhausted

Mitigation: Restarted database connection pool

Current status: Service partially restored, monitoring

Next update: 15:00 UTC or when major progress made

## Recovery Announcement

[Resolved] HashInsight Service Restored

Time: 2025-10-03 15:12 UTC

Duration: 49 minutes

Root cause: Improper PostgreSQL connection pool configuration

Solution: Increased connection pool limits and optimized slow queries

All services now fully restored.

Post-incident review will be published within 24 hours.

Thank you for your patience.

## 10.7 Post-Incident Review (Postmortem)

### Template

```
# HashInsight Incident Report - 2025-10-03 Database Outage
```

#### ## Summary

- \*\*Incident ID\*\*: INC-2025-1003-001
- \*\*Severity\*\*: P0
- \*\*Start Time\*\*: 2025-10-03 14:23 UTC
- \*\*Recovery Time\*\*: 2025-10-03 15:12 UTC
- \*\*Duration\*\*: 49 minutes
- \*\*Affected Users\*\*: 100% (all users)

#### ## Timeline

- 14:23 - Prometheus alert: Database connection failure
- 14:25 - On-call engineer confirms incident

- 14:30 - Identified connection pool exhaustion
- 14:35 - Implemented temporary mitigation (killed idle connections)
- 14:45 - Service partially restored
- 15:00 - Implemented permanent fix (increased connection pool)
- 15:12 - Service fully restored

#### ## Root Cause

PostgreSQL connection pool size configured at 10, but actual concurrent demand reached 30+, causing connection wait timeouts.

#### ## Resolution

1. Temporary: Killed idle connections, freed connection pool
2. Permanent: Increased connection pool size to 20, overflow to 40
3. Optimization: Identified and optimized 3 slow queries

#### ## Prevention Measures

- [ ] Set connection pool alerts (usage >80%)
- [ ] Regular slow query review
- [ ] Add capacity planning process
- [ ] Add connection pool monitoring dashboard

#### ## Lessons Learned

Quick response and good communication

Lacked connection pool monitoring

Insufficient capacity planning

## Appendices

---

### Appendix A: Quick Command Reference

#### Application Management

```
# Start application
gunicorn --bind 0.0.0.0:5000 --workers 4 main:app

# Graceful restart
kill -HUP $(cat /var/run/gunicorn.pid)

# Stop application
kill -TERM $(cat /var/run/gunicorn.pid)

# View processes
```

```
ps aux | grep gunicorn  
  
# View logs  
tail -f /var/log/hashinsight/app.log
```

## Database

```
# Connect to database  
psql $DATABASE_URL  
  
# View tables  
\dt  
  
# View connection count  
SELECT count(*) FROM pg_stat_activity;  
  
# Backup database  
pg_dump $DATABASE_URL > backup.sql  
  
# Restore database  
psql $DATABASE_URL < backup.sql
```

## Cache

```
# Redis connection  
redis-cli  
  
# View all keys  
redis-cli KEYS 'hashinsight:'  
  
# Clear cache  
redis-cli FLUSHDB  
  
# View memory usage  
redis-cli INFO memory
```

## Monitoring

```
# View health status  
curl http://localhost:5000/health | jq  
  
# View Prometheus metrics
```

```
curl http://localhost:9090/metrics  
  
# View SLO status  
curl http://localhost:5000/api/slo/status | jq
```

## Appendix B: Configuration File Examples

### .env Example

```
# Required configuration  
DATABASE_URL=postgresql://user:pass@host:5432/hashinsight_db  
SESSION_SECRET=your-secret-key-min-32-chars  
ENCRYPTION_PASSWORD=encryption-key-min-32-chars  
  
# Blockchain configuration  
BLOCKCHAIN_ENABLED=true  
BLOCKCHAIN_PRIVATE_KEY=0x1234567890abcdef...  
BLOCKCHAIN_NETWORK=base-sepolia  
BASE_RPC_URL=https://sepolia.base.org  
  
# Backup configuration  
BACKUP_DIR=/var/backups/hashinsight  
BACKUP_ENCRYPTION_KEY=backup-encryption-key  
BACKUP_RETENTION_DAYS=30  
BACKUP_STORAGE_TYPE=s3  
BACKUP_STORAGE_BUCKET=hashinsight-backups  
  
# KMS configuration (AWS)  
AWS_KMS_KEY_ID=arn:aws:kms:us-east-1:123456789:key/xxxxxx  
AWS_KMS_REGION=us-east-1  
AWS_ACCESS_KEY_ID=AKIAXXXXXXXX  
AWS_SECRET_ACCESS_KEY=xxxxxxxxxxxx  
  
# Monitoring configuration  
ENABLE_BACKGROUND_SERVICES=0 # Default: disabled; set to 1 to enable background data collection  
PROMETHEUS_PORT=9090  
SLO_MEASUREMENT_WINDOW=30  
  
# API configuration  
COINWARZ_API_KEY=your-coinwarz-api-key  
COINGECKO_API_KEY=your-coingecko-api-key
```

## systemd Service File Example

```
# /etc/systemd/system/hashinsight.service

[Unit]
Description=HashInsight Platform
After=network.target postgresql.service

[Service]
Type=notify
User=hashinsight
Group=hashinsight
WorkingDirectory=/opt/hashinsight
EnvironmentFile=/opt/hashinsight/.env

ExecStart=/opt/hashinsight/venv/bin/gunicorn \
--bind 0.0.0.0:5000 \
--workers 4 \
--threads 2 \
--worker-class gthread \
--timeout 120 \
--max-requests 1000 \
--access-logfile /var/log/hashinsight/access.log \
--error-logfile /var/log/hashinsight/error.log \
--pid /var/run/gunicorn.pid \
main:app

ExecReload=/bin/kill -s HUP $MAINPID
KillMode=mixed
KillSignal=SIGHTERM
TimeoutStopSec=30

Restart=on-failure
RestartSec=10

[Install]
WantedBy=multi-user.target
```

## Appendix C: Monitoring Metrics Reference

### Prometheus Metrics List

Metric Name	Type	Description
hashinsight_requests_total	Counter	Total request count
hashinsight_request_latency_seconds	Histogram	Request latency
hashinsight_error_rate	Gauge	Error rate
hashinsight_cache_hit_rate	Gauge	Cache hit rate
hashinsight_db_query_duration_seconds	Histogram	Database query duration
hashinsight_db_connection_pool_size	Gauge	Database connection pool size
hashinsight_db_connection_pool_active	Gauge	Active connection count
hashinsight_slo_compliance	Gauge	SLO compliance
hashinsight_slo_error_budget_remaining	Gauge	Error budget remaining
hashinsight_circuit_breaker_state	Gauge	Circuit breaker state

### SLO Thresholds

SLO Type	Target	Warning Threshold	Critical Threshold
Availability	99.95%	<99.9%	<99.5%
P95 latency	≤250ms	>200ms	>300ms
Error rate	≤0.1%	>0.5%	>1%
Error budget	21.6min/month	<20%	<10%

## Appendix D: Error Code Reference

Error Code	HTTP Status	Description	Solution
AUTH_001	401	Not logged in	Re-login
AUTH_002	403	Insufficient permissions	Contact administrator
AUTH_003	401	Session expired	Re-login
DB_001	500	Database connection failed	Check DATABASE_URL
DB_002	500	Query timeout	Optimize query or increase timeout
CACHE_001	503	Redis unavailable	Check Redis service
API_001	429	Too many requests	Reduce request frequency
API_002	500	External API failed	Retry later
BLOCKCHAIN_001	500	Private key not configured	Set BLOCKCHAIN_PRIVATE_KEY
ENCRYPTION_001	500	Encryption key not configured	Set ENCRYPTION_PASSWORD

## Appendix E: Version History

Version	Date	Author	Changes
v2.0	2025-10-03	HashInsight Ops Team	Complete operations manual first release
v1.9	2025-09-15	DevOps Team	Added Request Coalescing optimization
v1.8	2025-08-20	Security Team	Enhanced KMS and mTLS documentation
v1.7	2025-07-10	Platform Team	Added SLO monitoring chapter

---

## Document Maintenance

---

**Maintenance Responsibility:** HashInsight Platform Operations Team

**Review Cycle:** Quarterly

**Feedback Channel:** ops@hashinsight.net

**Documentation Repository:** <https://github.com/hashinsight/operations-manual>

**Last Updated:** October 3, 2025

**Next Review:** January 3, 2026

---

**© 2025 HashInsight Platform. All Rights Reserved.**

**Confidential - Internal Use Only**